

Probabilistic Programming for Planning as Inference Problems

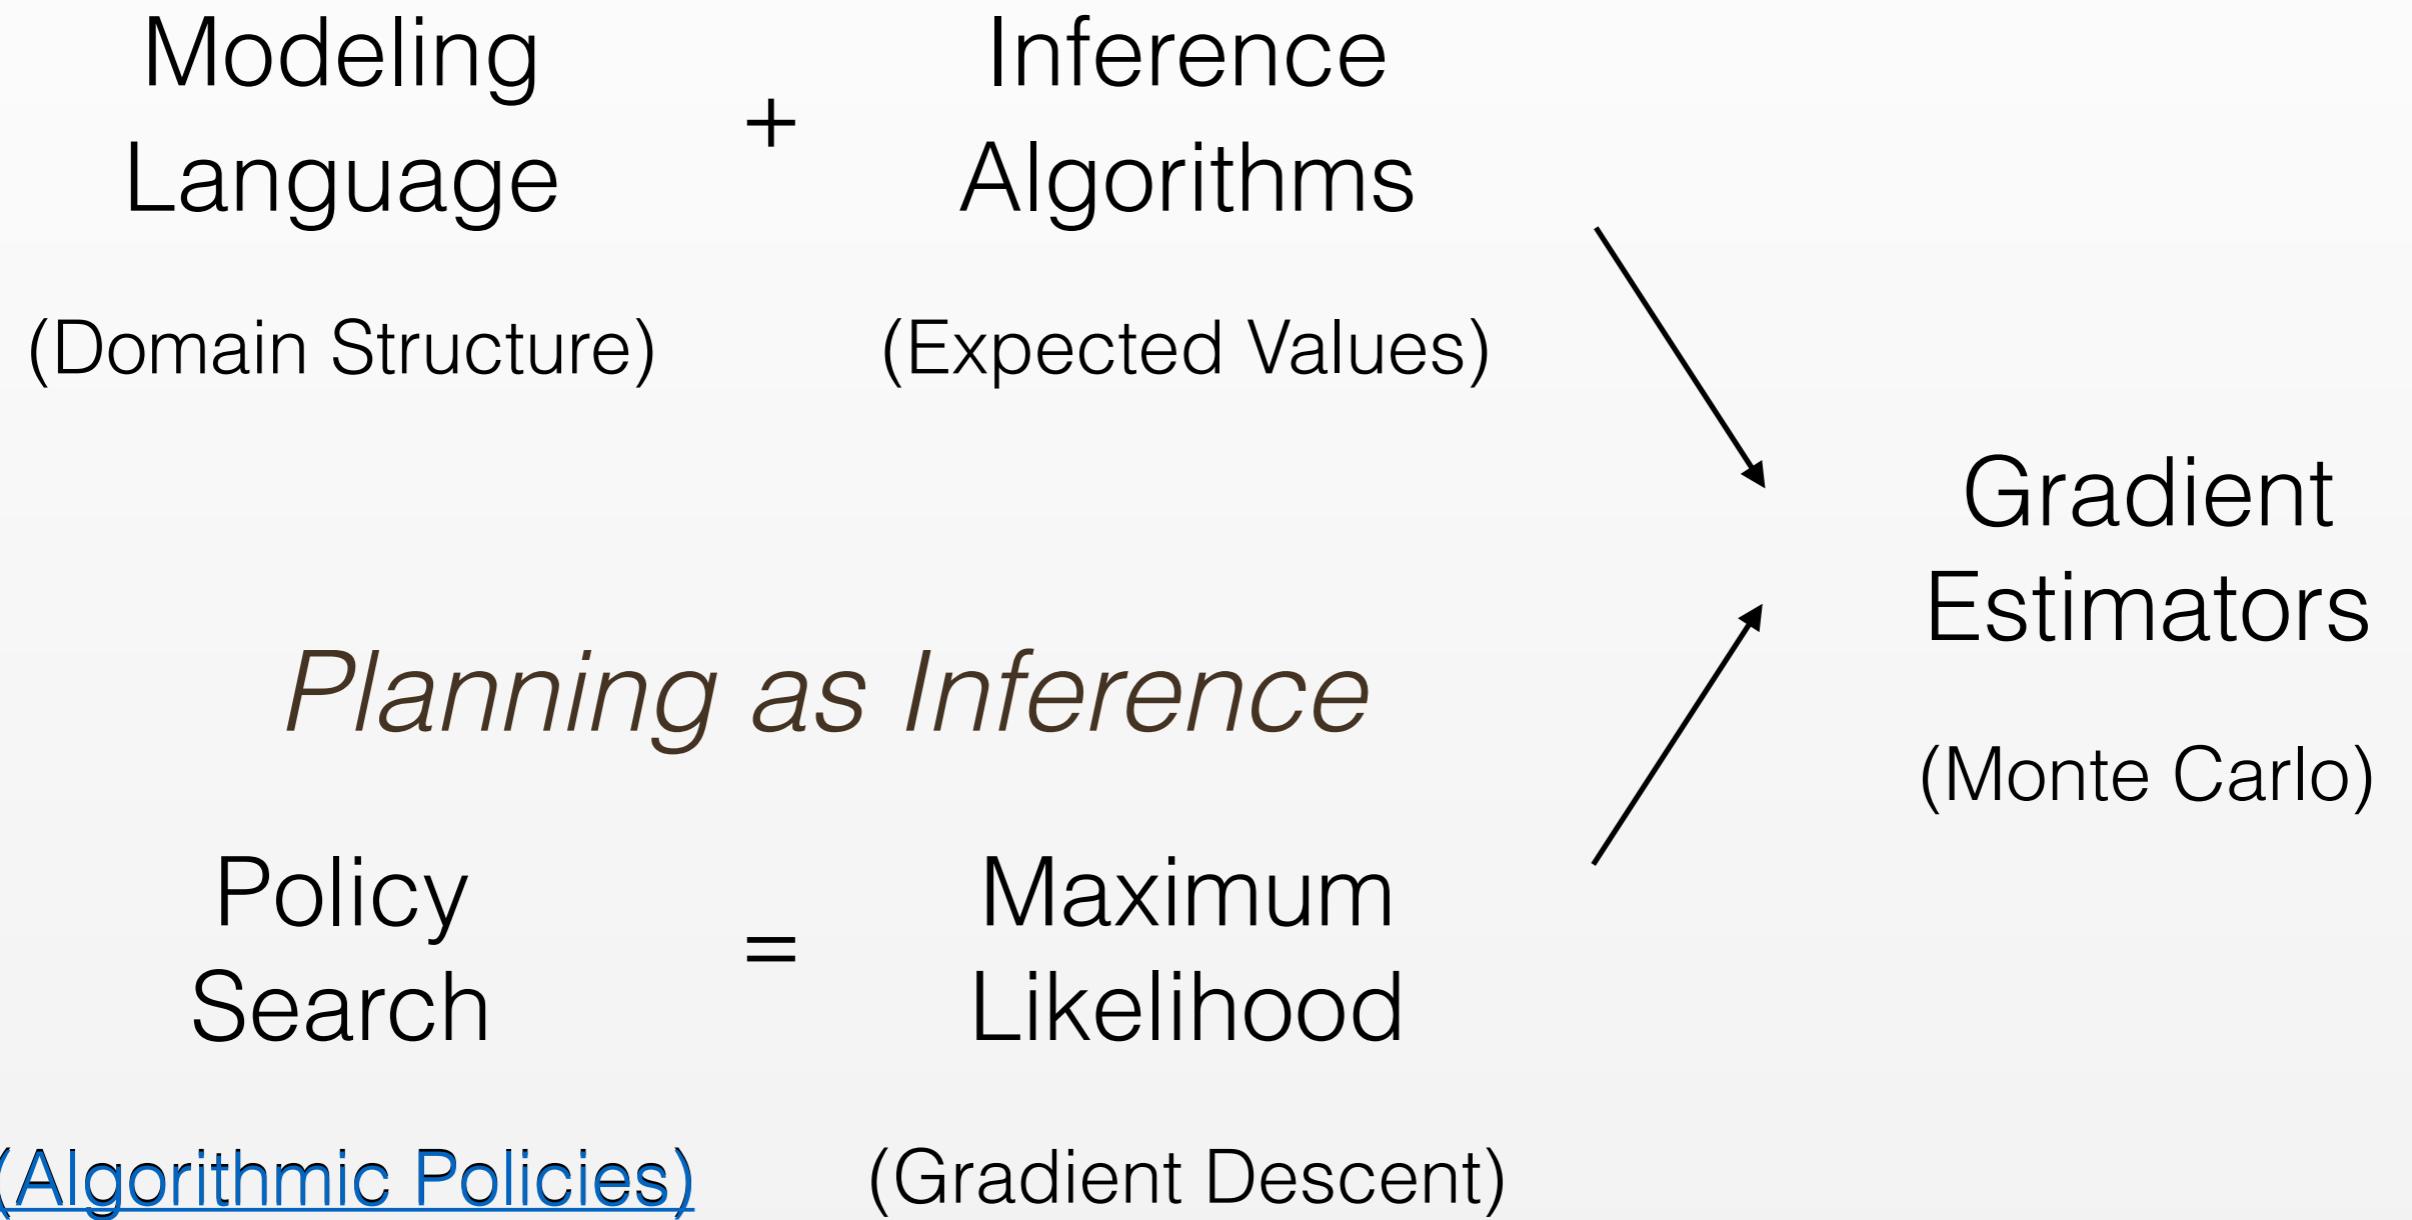
Jan-Willem van de Meent, Northeastern University



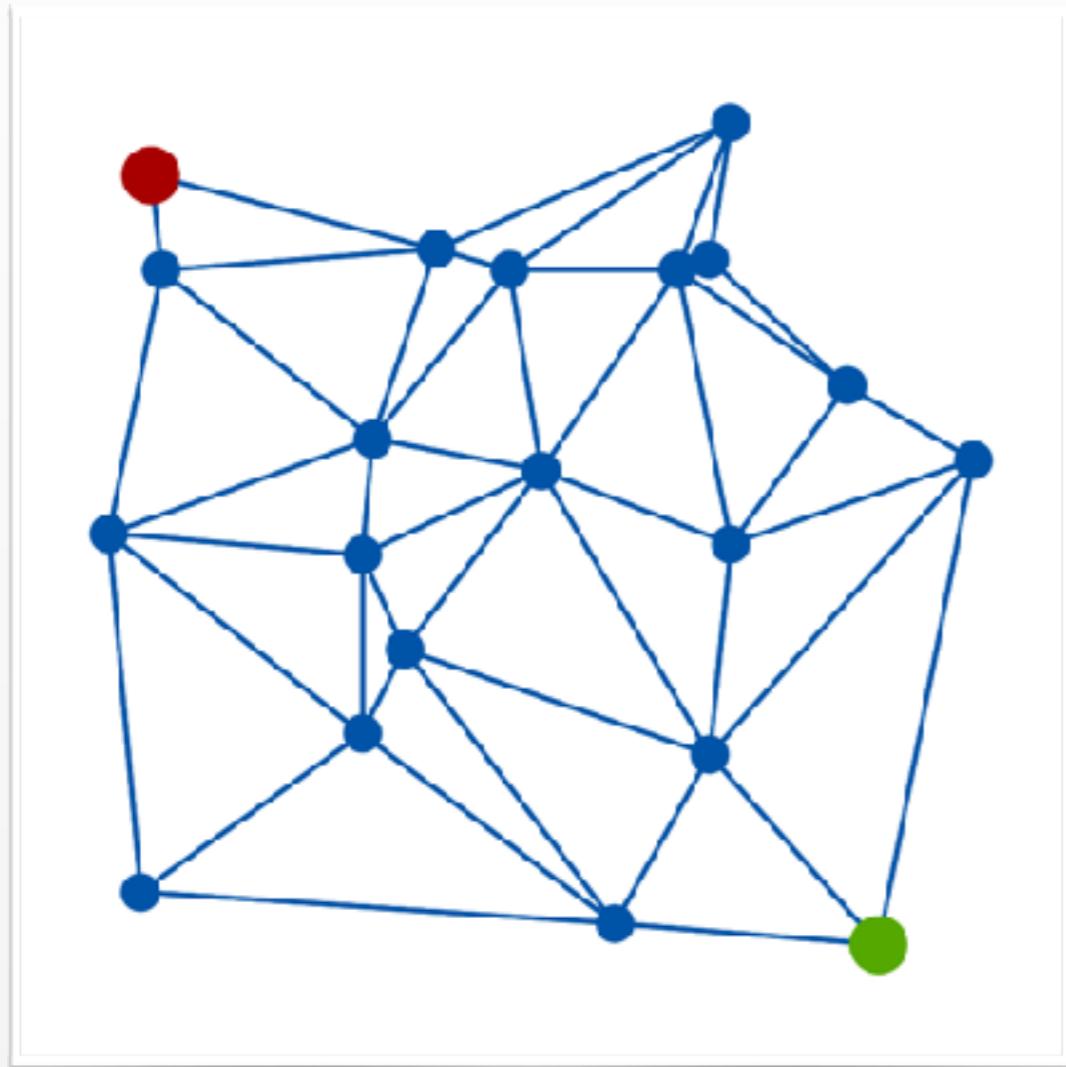
[van de Meent, Paige, Tolpin, Wood, AISTATS 2016]

[Siddharth, Paige, van de Meent, Desmaison, Goodman, Kohli, Wood, Torr, NIPS 2017]

Probabilistic Programming

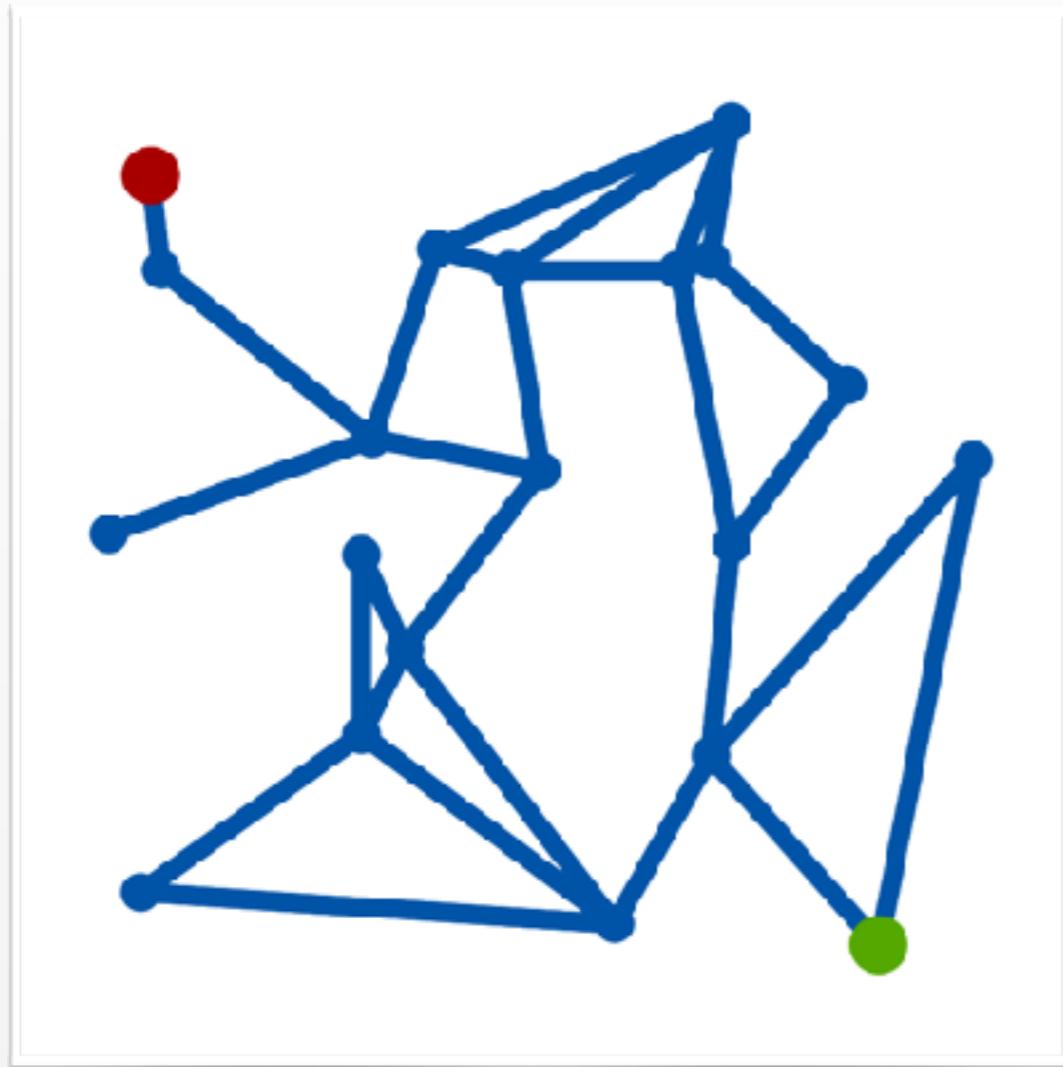


Canadian Traveler Problem



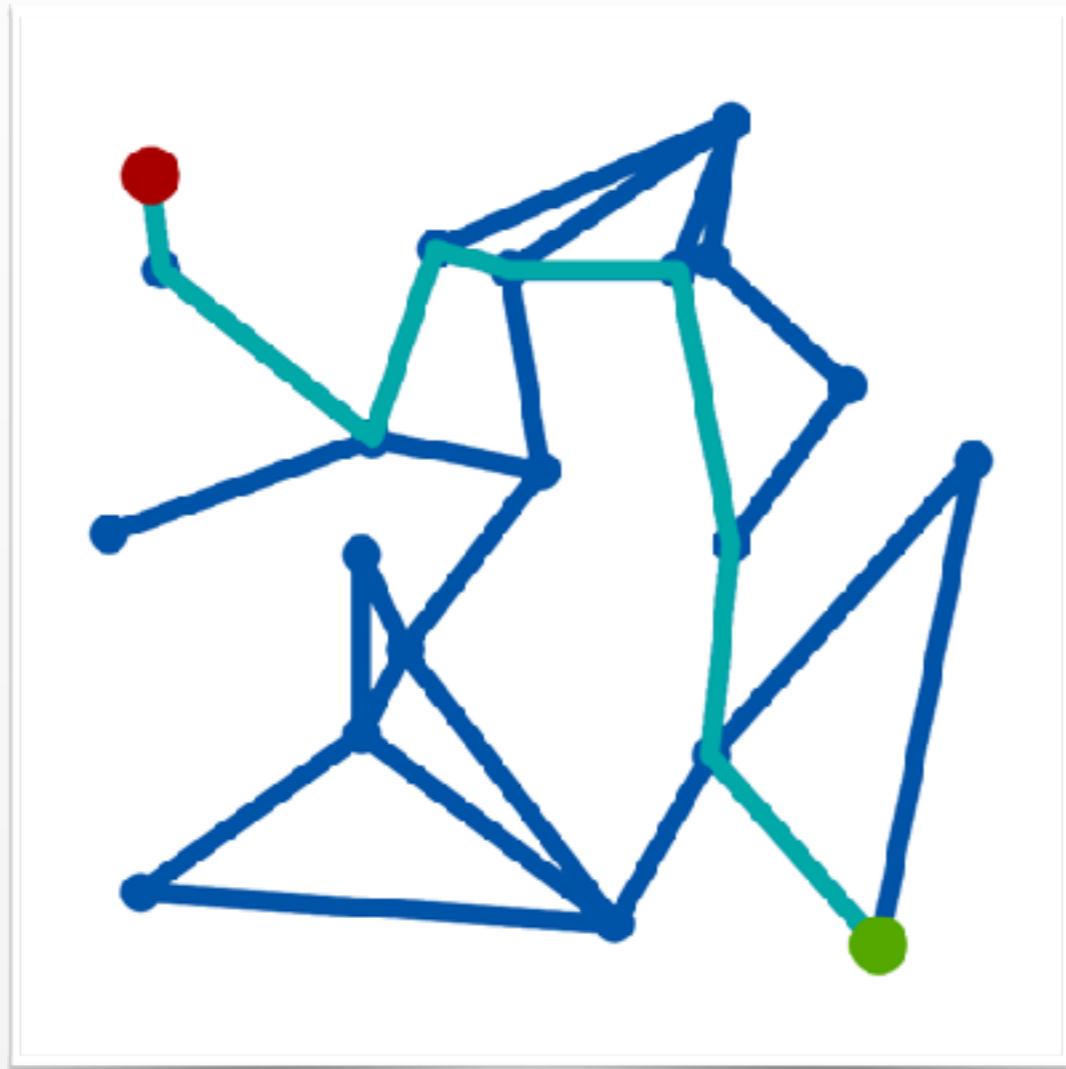
Known network of roads

Canadian Traveler Problem



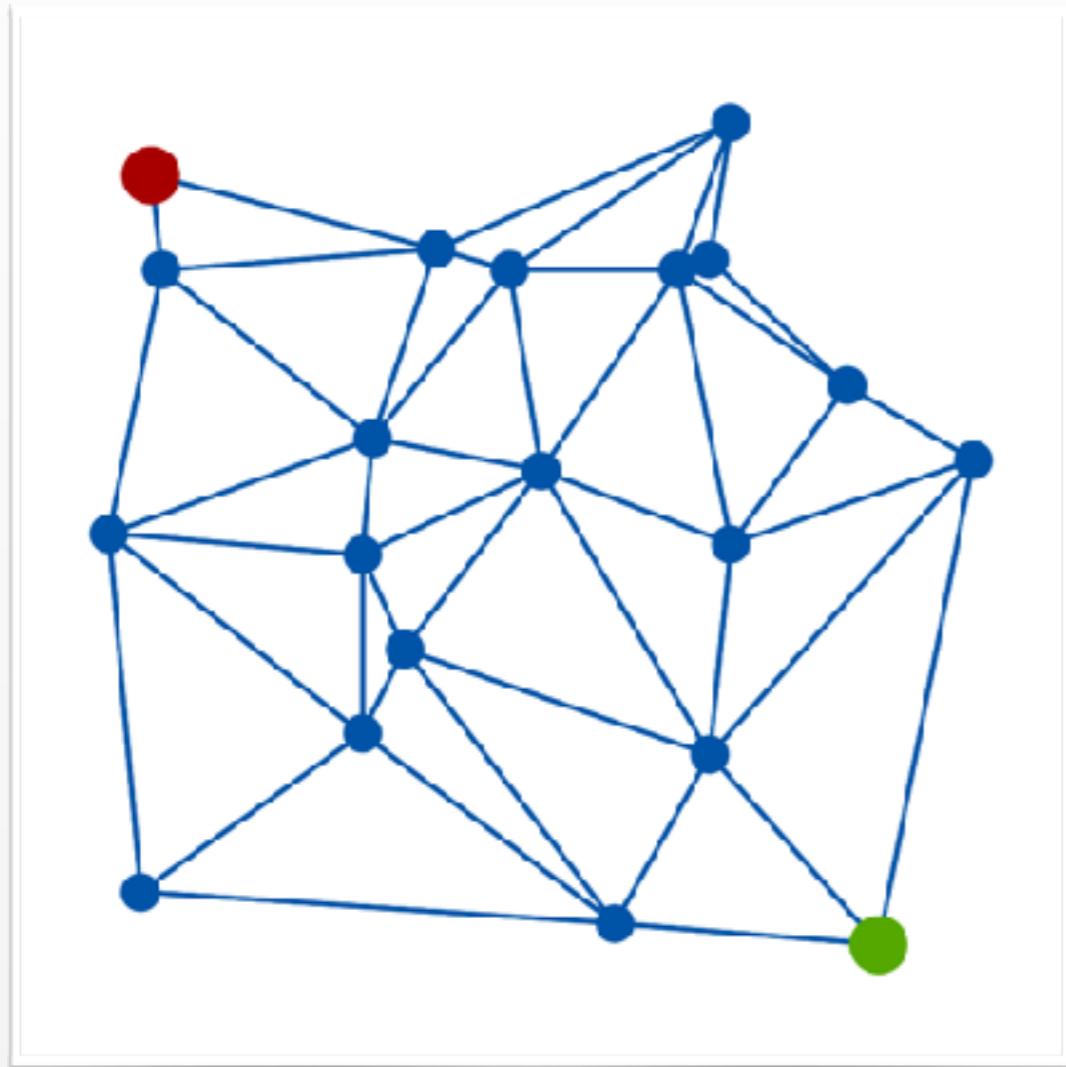
Edges missing at Random

Canadian Traveler Problem



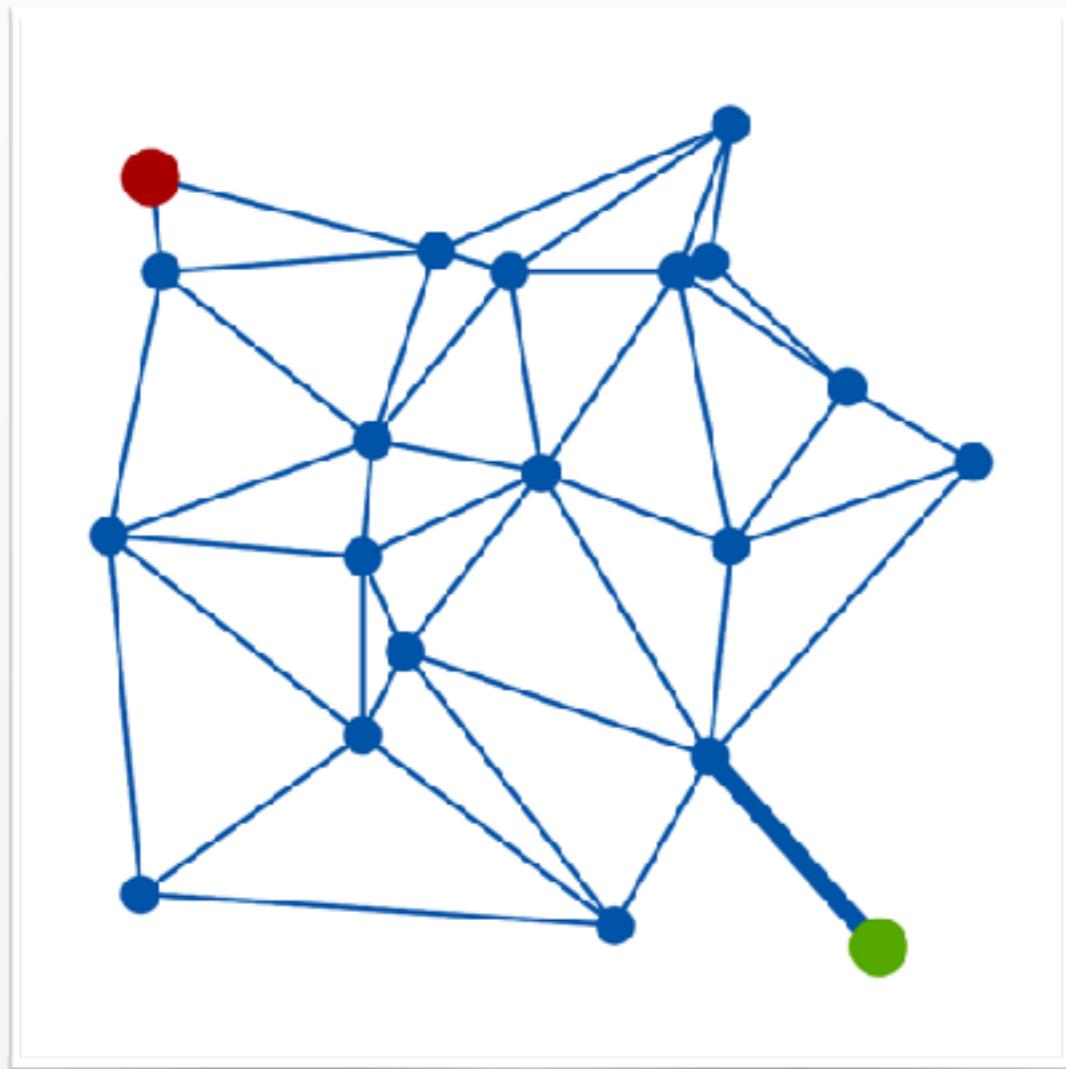
Perfect information:
Dijkstra's algorithm

Canadian Traveler Problem



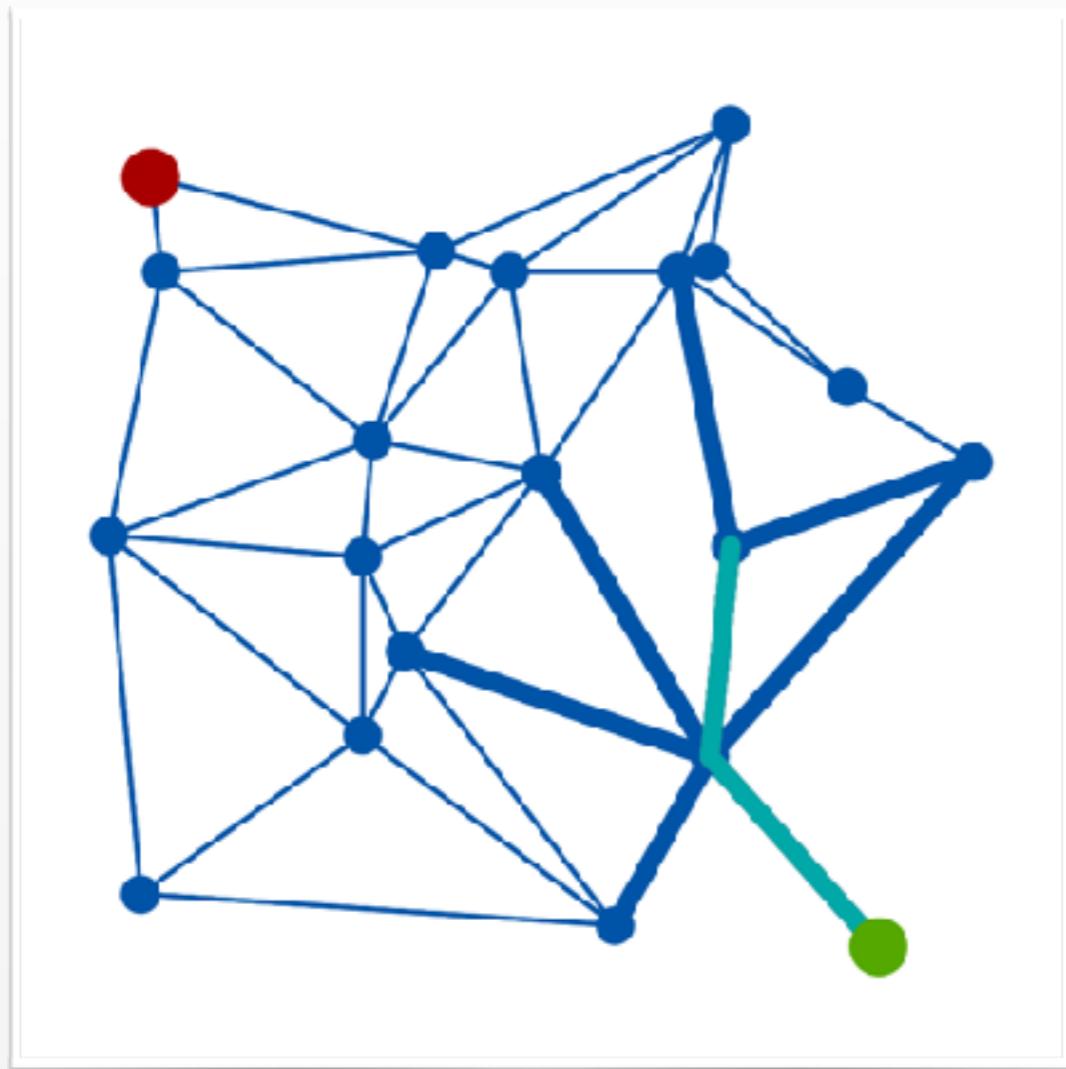
State of roads unknown:
Planning problem

Canadian Traveler Problem



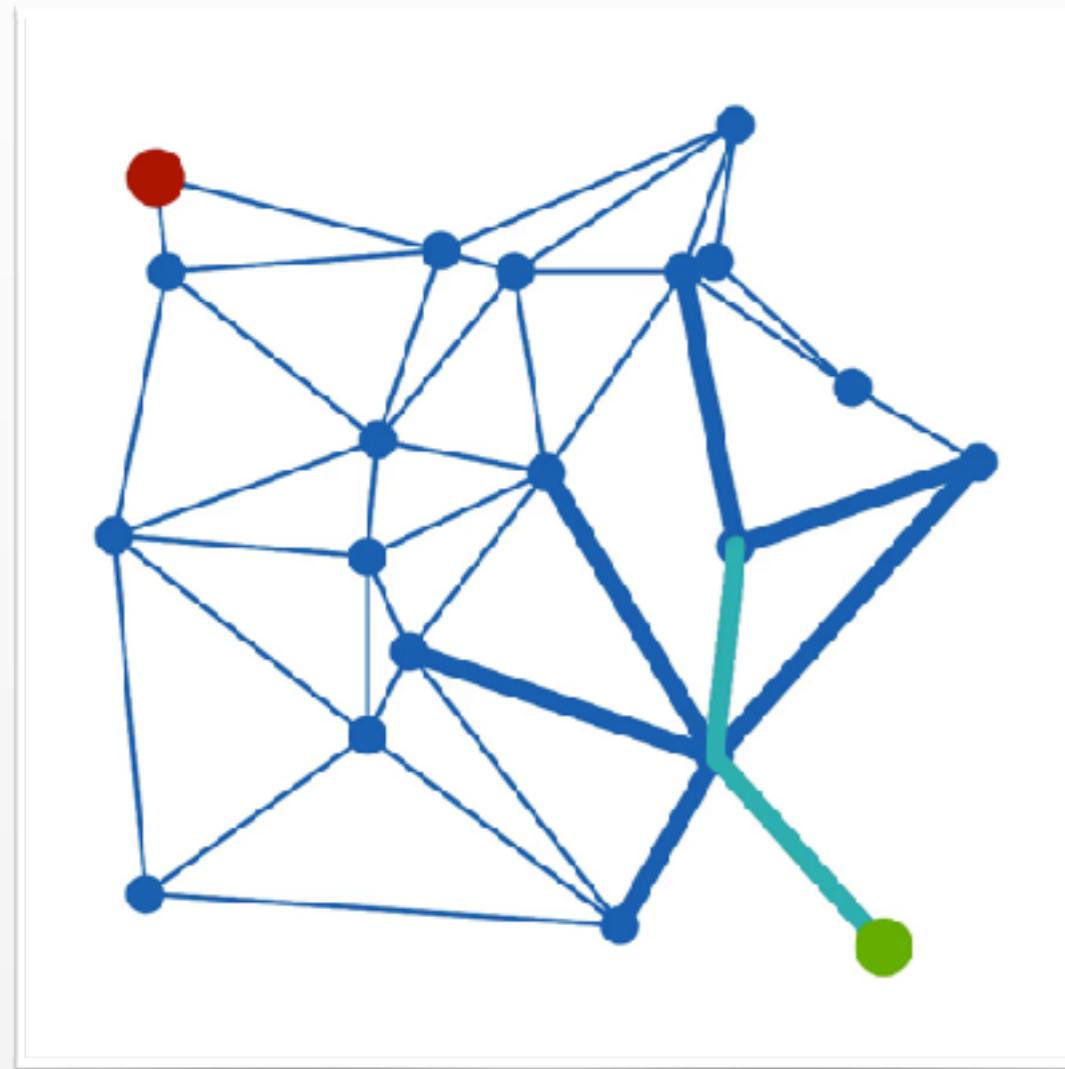
State of roads unknown:
Planning problem

Canadian Traveler Problem



State of roads unknown:
Planning problem

Canadian Traveler Problem



State of roads unknown:
Planning problem

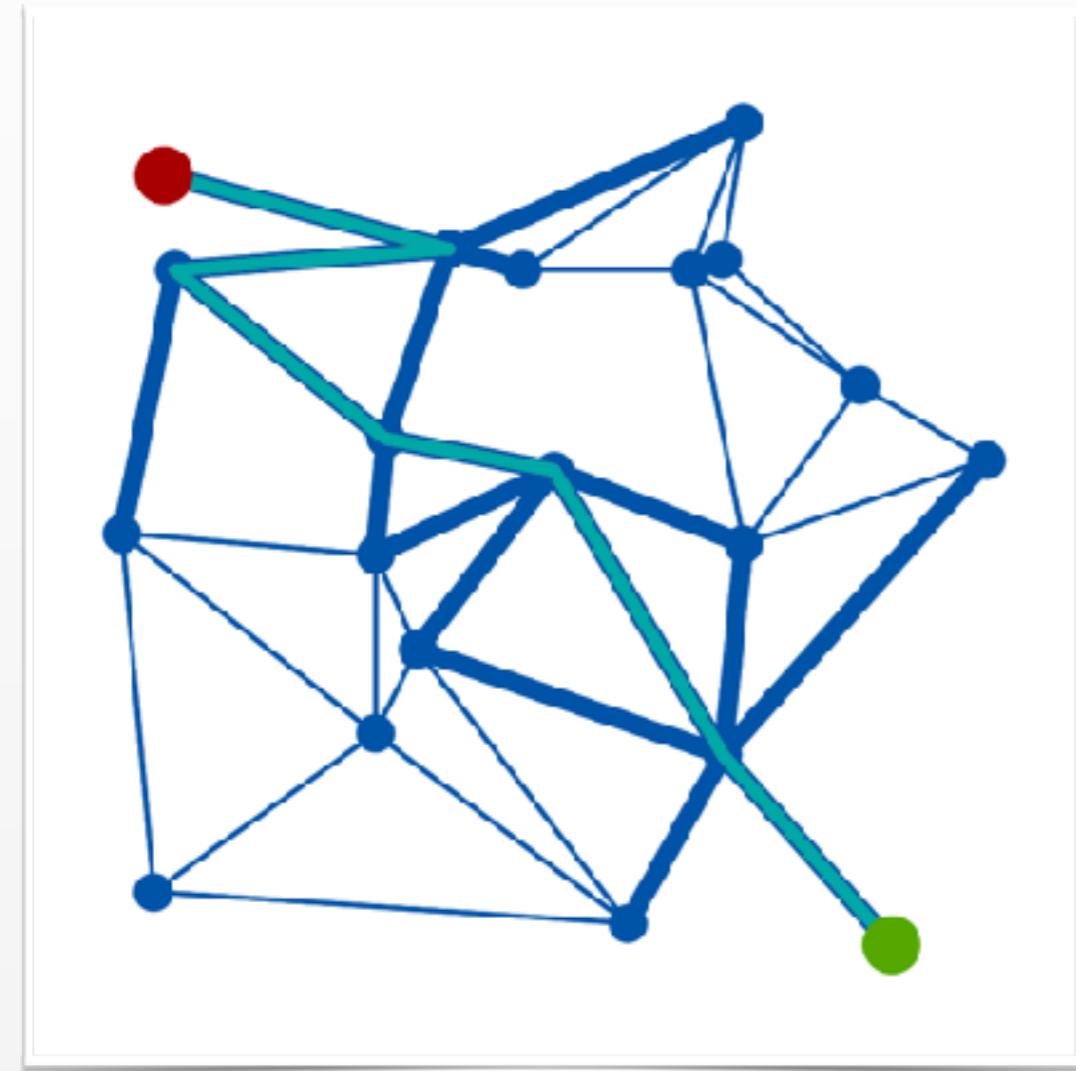
Strawman Policy

- Perform Depth First Search
- Choose unexplored paths uniformly at random
- Guaranteed to reach goal (if accessible)

Planning with Prob. Prog.

- Express policies as probabilistic programs
- Optimize free parameters to maximize expected reward

Policy Search (upper-level)



State of roads unknown:
Planning problem

Policy (Probabilistic Program)

$$\theta \sim p(\theta; \lambda)$$

$$a_t \sim p(a_t | s_t, \theta)$$

World (Probabilistic Program)

$$s_t \sim p(s_t | s_{t-1}, a_t)$$

Maximize Expected Reward

$$x := (a_0, s_1, \dots, s_T, a_T)$$

$$R(x) := \sum_t -\text{distance}(a_t)$$

$$J(\lambda) := \mathbb{E}_{p(x|\lambda)} [R(x)]$$

Probabilistic Programs for Inference

Probabilistic Program

```
(defquery lin-reg [x-vals y-vals]
  (let [a (sample (normal 0 1))
        b (sample (normal 0 1))
        f (fn [x] (+ (* a x) b))]
    (map (fn [x y]
           (observe
             (normal (f x) 0.1) y))
         x-vals y-vals)
    [a b]))
```

Joint Probability

$$p(\mathbf{y}, \mathbf{x}) = p(\mathbf{y} | \mathbf{x}) p(\mathbf{x})$$

Likelihood Prior

Special Forms

sample random value \mathbf{x}

observe condition on value \mathbf{y}

Probabilistic Programs for Inference

Probabilistic Program

```
(defquery lin-reg [x-vals y-vals]
  (let [a (sample (normal 0 1))
        b (sample (normal 0 1))
        f (fn [x] (+ (* a x) b))]
    (map (fn [x y]
           (observe
             (normal (f x) 0.1) y))
         x-vals y-vals)
    [a b]))
```

Posterior Density

$$p(\mathbf{x} | \mathbf{y}) = p(\mathbf{y} | \mathbf{x})p(\mathbf{x})/p(\mathbf{y})$$

*Distribution on **sample** values given **observe** values.*

Special Forms

sample random value \mathbf{x}

observe condition on value \mathbf{y}

Probabilistic Programs for Inference

Probabilistic Program

```
(defquery lin-reg [x-vals y-vals]
  (let [a (sample (normal 0 1))
        b (sample (normal 0 1))
        f (fn [x] (+ (* a x) b))]
    (map (fn [x y]
           (observe
            (normal (f x) 0.1) y))
         x-vals y-vals)
    [a b]))
```

Special Forms

sample random value \mathbf{x}

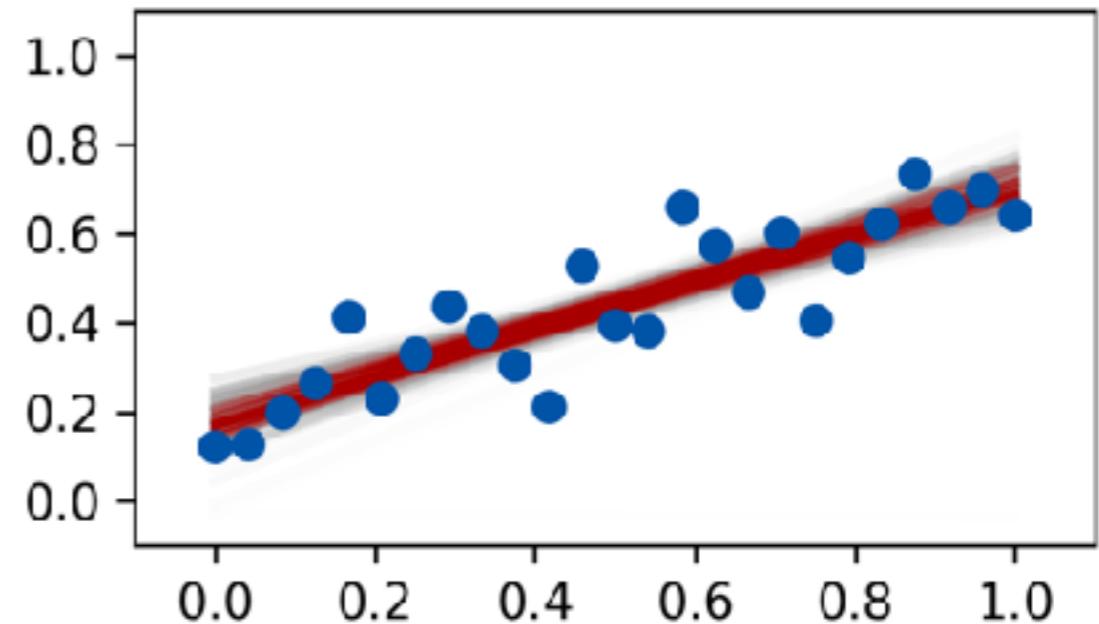
observe condition on value \mathbf{y}

Posterior Density

$$p(\mathbf{x} | \mathbf{y}) = p(\mathbf{y} | \mathbf{x})p(\mathbf{x})/p(\mathbf{y})$$

Distribution on **sample** values given **observe** values.

Inference



Probabilistic Programs for Planning

Probabilistic Program

```
(let [θ (sample (policy-prior λ))
      P (make-policy θ)
      x (sample-episode P)]
  (factor (reward x)))
[x θ])
```

Special Forms

sample random value \mathbf{x}

factor assign log probability

Conditioning on Reward

$\exp \mathbf{R}(\mathbf{x})$ replaces $p(\mathbf{y}|\mathbf{x})$

Policy Search

Maximize expected **reward**,
marginalizing over trajectories \mathbf{x} ,
with respect to hyperparameters λ

Bayesian Analogue: Maximum Likelihood, Empirical Bayes

Planning as Inference

Target
Density

$$\begin{aligned}\pi(\mathbf{x}; \boldsymbol{\phi}) &:= \gamma(\mathbf{x}; \boldsymbol{\phi})/Z \\ Z &:= \int \gamma(\mathbf{x}; \boldsymbol{\phi}) d\mathbf{x} \quad (\text{intractable integral})\end{aligned}$$

Bayesian
Inference

$$\begin{aligned}\gamma(\mathbf{x}; \boldsymbol{\phi}) &:= p(\mathbf{y}, \mathbf{x}; \boldsymbol{\phi}) \\ Z &:= \int p(\mathbf{y}, \mathbf{x}; \boldsymbol{\phi}) d\mathbf{x} = p(\mathbf{y}; \boldsymbol{\phi})\end{aligned}$$

Planning

$$\begin{aligned}\gamma(\mathbf{x}; \boldsymbol{\phi}) &:= \begin{cases} p(\mathbf{x}; \boldsymbol{\phi}) R(\mathbf{x}) & R(\mathbf{x}) \geq 0 \\ p(\mathbf{x}; \boldsymbol{\phi}) \exp[R(\mathbf{x})] & R(\mathbf{x}) \in \mathbb{R} \end{cases} \\ Z &:= \begin{cases} \mathbb{E}_{p(\mathbf{x}; \boldsymbol{\phi})}[R(\mathbf{x})] & \text{Expected reward} \\ \mathbb{E}_{p(\mathbf{x}; \boldsymbol{\phi})} \exp[R(\mathbf{x})] & \text{Desirability} \end{cases}\end{aligned}$$

Optimizing Lower Bounds

$$Z(\phi) := \int \gamma(\mathbf{x}; \phi) d\mathbf{x} = \int q(\mathbf{x}; \lambda) \frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} d\mathbf{x} = \mathbb{E}_{q(\mathbf{x}; \lambda)} \left[\frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right]$$

Approximation: Lower Bound (from Jensen's Inequality)

$$\mathcal{L}(\phi, \lambda) := \mathbb{E}_{q(\mathbf{x}; \lambda)} \left[\log \frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right] \leq \log \mathbb{E}_{q(\mathbf{x}; \lambda)} \left[\frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right] = \log Z(\phi)$$

Variational Inference: $\max_{\lambda} \mathcal{L}(\phi, \lambda) = \min_{\lambda} \text{KL}(q(\mathbf{x}; \lambda) \mid \pi(\mathbf{x}; \phi))$

Maximum Likelihood: $\max_{\phi} \mathcal{L}(\phi, \lambda) = \max_{\phi} \log Z(\phi)$

Optimizing Lower Bounds

$$\mathcal{L}(\boldsymbol{\phi}, \boldsymbol{\lambda}) := \mathbb{E}_{q(\mathbf{x}; \boldsymbol{\lambda})} \left[\log \frac{\gamma(\mathbf{x}; \boldsymbol{\phi})}{q(\mathbf{x}; \boldsymbol{\lambda})} \right] \leq \log Z$$

Policy Search

$$\gamma(\mathbf{x}; \boldsymbol{\phi}) = p(\mathbf{x}; \boldsymbol{\phi}) \exp[\mathcal{R}(\mathbf{x})] \quad \mathcal{L}(\boldsymbol{\lambda}, \boldsymbol{\lambda}) = \mathbb{E}_{p(\mathbf{x}; \boldsymbol{\lambda})} \left[\log \frac{p(\mathbf{x}, \boldsymbol{\lambda}) \exp \mathcal{R}(\mathbf{x})}{p(\mathbf{x}, \boldsymbol{\lambda})} \right]$$

$$q(\mathbf{x}; \boldsymbol{\lambda}) = p(\mathbf{x}; \boldsymbol{\lambda}) \quad = \mathbb{E}_{p(\mathbf{x}; \boldsymbol{\lambda})} [\mathcal{R}(\mathbf{x})]$$

$$Z = \mathbb{E}_{p(\mathbf{x}; \boldsymbol{\phi})} [\exp \mathcal{R}(\mathbf{x})] \quad = J(\boldsymbol{\lambda})$$

Core Problem: Gradient Estimation

Easy: Gradient with respect to generative parameters ϕ

$$\nabla_{\phi} \mathcal{L}(\phi, \lambda) := \nabla_{\phi} \mathbb{E}_{q(\mathbf{x}; \lambda)} \left[\log \frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right] = \mathbb{E}_{q(\mathbf{x}; \lambda)} [\nabla_{\phi} \log \gamma(\mathbf{x}; \phi)]$$

$$\approx \frac{1}{S} \sum_{s=1}^S \nabla_{\phi} \log \gamma(\mathbf{x}^{(s)}; \phi) \quad \mathbf{x}^{(s)} \sim q(\mathbf{x}; \lambda)$$

Problem: generative model $\gamma(\mathbf{x}; \phi)$ must be differentiable
(*not generally true for probabilistic programs*)

Trick 1: Likelihood-ratio Estimators

Harder: Gradient w.r.t generative parameters λ

$$\nabla_{\lambda} \mathcal{L}(\phi, \lambda) := \nabla_{\lambda} \mathbb{E}_{q(\mathbf{x}; \lambda)} \left[\log \frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right]$$

$$= \underbrace{\int \nabla_{\lambda} q(\mathbf{x}; \lambda) \log \left[\frac{\gamma(\mathbf{x}; \phi)}{q(\mathbf{x}; \lambda)} \right] - q(\mathbf{x}; \lambda) \nabla_{\lambda} \log q(\mathbf{x}; \lambda) d\mathbf{x}}$$

Trick 1 (REINFORCE)

$$\nabla_{\lambda} q(\mathbf{x}; \lambda) = q(\mathbf{x}; \lambda) \nabla_{\lambda} \log q(\mathbf{x}; \lambda)$$

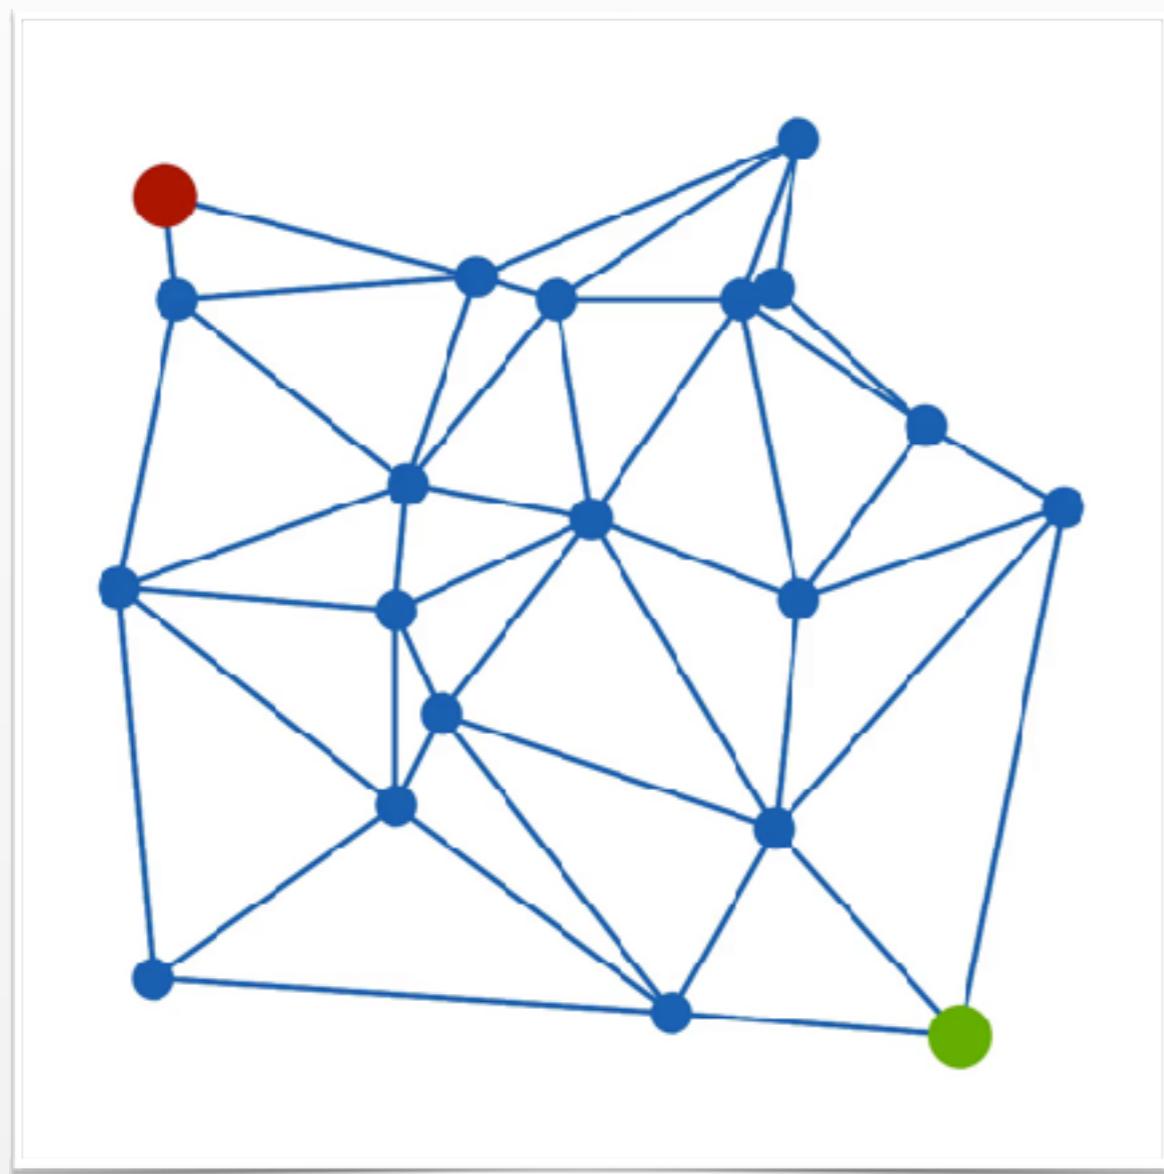
Requires grad
of $\log q(\mathbf{x}; \phi)$
but not $\gamma(\mathbf{x}; \phi)$

$$\approx \frac{1}{S} \sum_{s=1}^S \nabla_{\lambda} \log q(\mathbf{x}^{(s)}; \lambda) \left[\log \frac{\gamma(\mathbf{x}^{(s)}; \phi)}{q(\mathbf{x}^{(s)}; \lambda)} - \hat{b} \right] \quad \mathbf{x}^{(s)} \sim q(\mathbf{x}; \lambda)$$

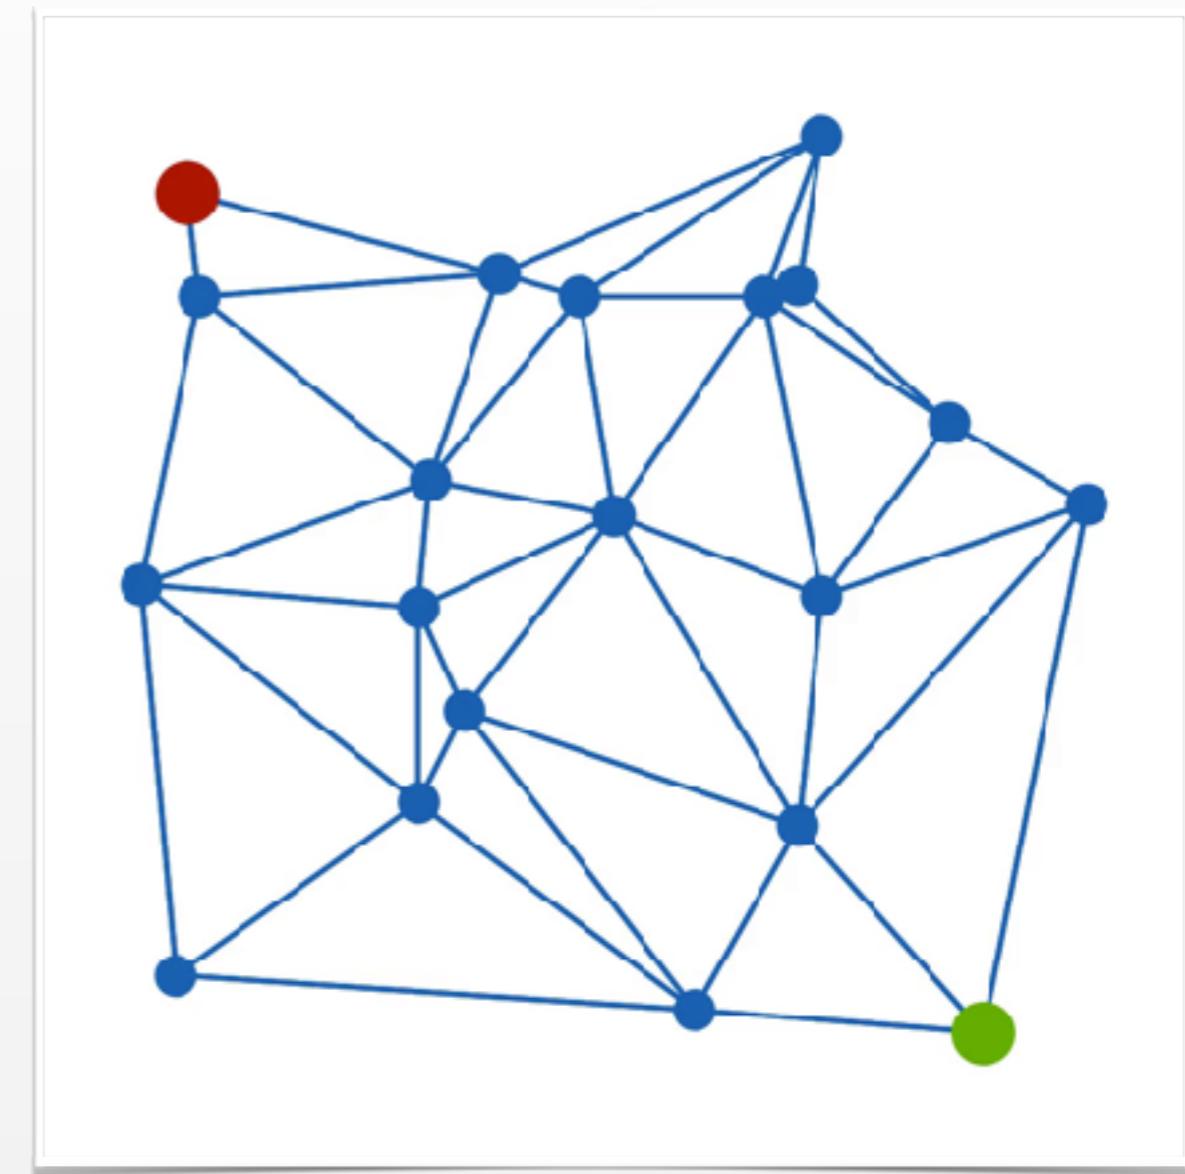
[see: Williams 1992, Wingate & Weber 2013, Ranganath et al. 2014]

Canadian Traveler Problem

Uniform Policy



Learned Policy (REINFORCE)



Limitation: Need 1000 samples per gradient step

Trick 2: Reparameterization

Idea: Define parameterized transformation $\mathbf{x} = \mathbf{X}(\boldsymbol{\epsilon}, \boldsymbol{\lambda})$

$$\left. \begin{array}{l} \boldsymbol{\epsilon} \sim p(\boldsymbol{\epsilon}) \\ \mathbf{x} = \mathbf{X}(\boldsymbol{\epsilon}; \boldsymbol{\lambda}) \end{array} \right\} \mathbf{x} \sim q(\mathbf{x}; \boldsymbol{\lambda}) \quad \left. \begin{array}{l} \boldsymbol{\epsilon} \sim \text{Normal}(\mathbf{0}, I) \\ \mathbf{x} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon} \end{array} \right\} \mathbf{x} \sim \text{Normal}(\boldsymbol{\mu}, \boldsymbol{\sigma} I)$$

Advantage: No longer need REINFORCE trick

$$\nabla_{\boldsymbol{\lambda}} \mathcal{L}(\boldsymbol{\phi}, \boldsymbol{\lambda}) := \nabla_{\boldsymbol{\lambda}} \mathbb{E}_{p(\boldsymbol{\epsilon})} \left[\log \frac{\gamma(\mathbf{X}(\boldsymbol{\epsilon}; \boldsymbol{\lambda}); \boldsymbol{\phi})}{q(\mathbf{X}(\boldsymbol{\epsilon}; \boldsymbol{\lambda}); \boldsymbol{\lambda})} \right] = \mathbb{E}_{p(\boldsymbol{\epsilon})} \left[\nabla_{\boldsymbol{\lambda}} \log \frac{\gamma(\mathbf{X}(\boldsymbol{\epsilon}; \boldsymbol{\lambda}); \boldsymbol{\phi})}{q(\mathbf{X}(\boldsymbol{\epsilon}; \boldsymbol{\lambda}); \boldsymbol{\lambda})} \right]$$

$$\simeq \underbrace{\frac{1}{S} \sum_{s=1}^S \nabla_{\boldsymbol{\lambda}} \log \frac{\gamma(\mathbf{X}(\boldsymbol{\epsilon}^{(s)}; \boldsymbol{\lambda}); \boldsymbol{\phi})}{q(\mathbf{X}(\boldsymbol{\epsilon}^{(s)}; \boldsymbol{\lambda}); \boldsymbol{\lambda})}}_{\text{Compute with AD}} \quad \boldsymbol{\epsilon}^{(s)} \sim p(\boldsymbol{\epsilon})$$

Deep Probabilistic Programs

Encoder $x \sim q(\mathbf{x}; \mathbf{y}, \boldsymbol{\lambda})$

```
class Encoder(torch.nn.Module):
    def __init__(self, x_sz, h_sz, y_sz, z_sz):
        # initializes layers: h, y_log_weights, ...
        ...

    def forward(self, x, y_values=None):
        q = probtorch.Trace()
        h = self.h(x)
        y = q.concrete(
            self.y_log_weights(h), 0.66,
            value=y_values, name='y')
        hy = torch.cat([h, y], -1)
        z = q.normal(self.z_mean(hy),
                     self.z_std(hy),
                     name='z')
        return q
```

Decoder $\gamma(\mathbf{x}; \mathbf{y}, \boldsymbol{\phi})$

```
class Decoder(torch.nn.Module):
    def __init__(self, x_sz, h_sz, y_sz, z_sz):
        # initializes layers: h, x_mean, ...
        ...

    def forward(self, x, q):
        p = probtorch.Trace()
        y = p.concrete(self.y_log_weights, 0.66,
                        value=q['y'], name='y')
        z = p.normal(0.0, 1.0,
                     value=q['z'], name='z')
        h = self.h(torch.cat([y, z], -1))
        x = p.loss(self.bce,
                   self.x_mean(h), x,
                   name='x')
        return p
```

Edward



Probabilistic Torch



Pyro



Gradient Estimation for Probabilistic Programs

		Generative Model $\gamma(\mathbf{x}; \mathbf{y}, \boldsymbol{\phi})$
Differentiable?	No	Yes
Approximating Distribution	No	Infer.NET
$q(\mathbf{x}; \mathbf{y}, \boldsymbol{\lambda})$	Yes	Relatively Unexplored ↑ ProbTorch, Edward, Pyro, Stan ↓ BBVI Inf. Comp ↔

Learning neural approximations [Ritchie et al. 2016, Le et al. 2017]

RELAX: REINFORCE + Differentiable surrogate [Grathwohl et al. 2017]

Algorithmic Proposals (e.g. incorporate MCTS)

Contributors

Anglican, Policy Search, Inference as Planning



David Tolpin



Brooks Paige



Hongseok Yang



Frank Wood



Tom Rainforth

Probabilistic Torch



Brooks
Paige



Siddharth N.



Alban
Desmaison



Alicant
Bozkurt

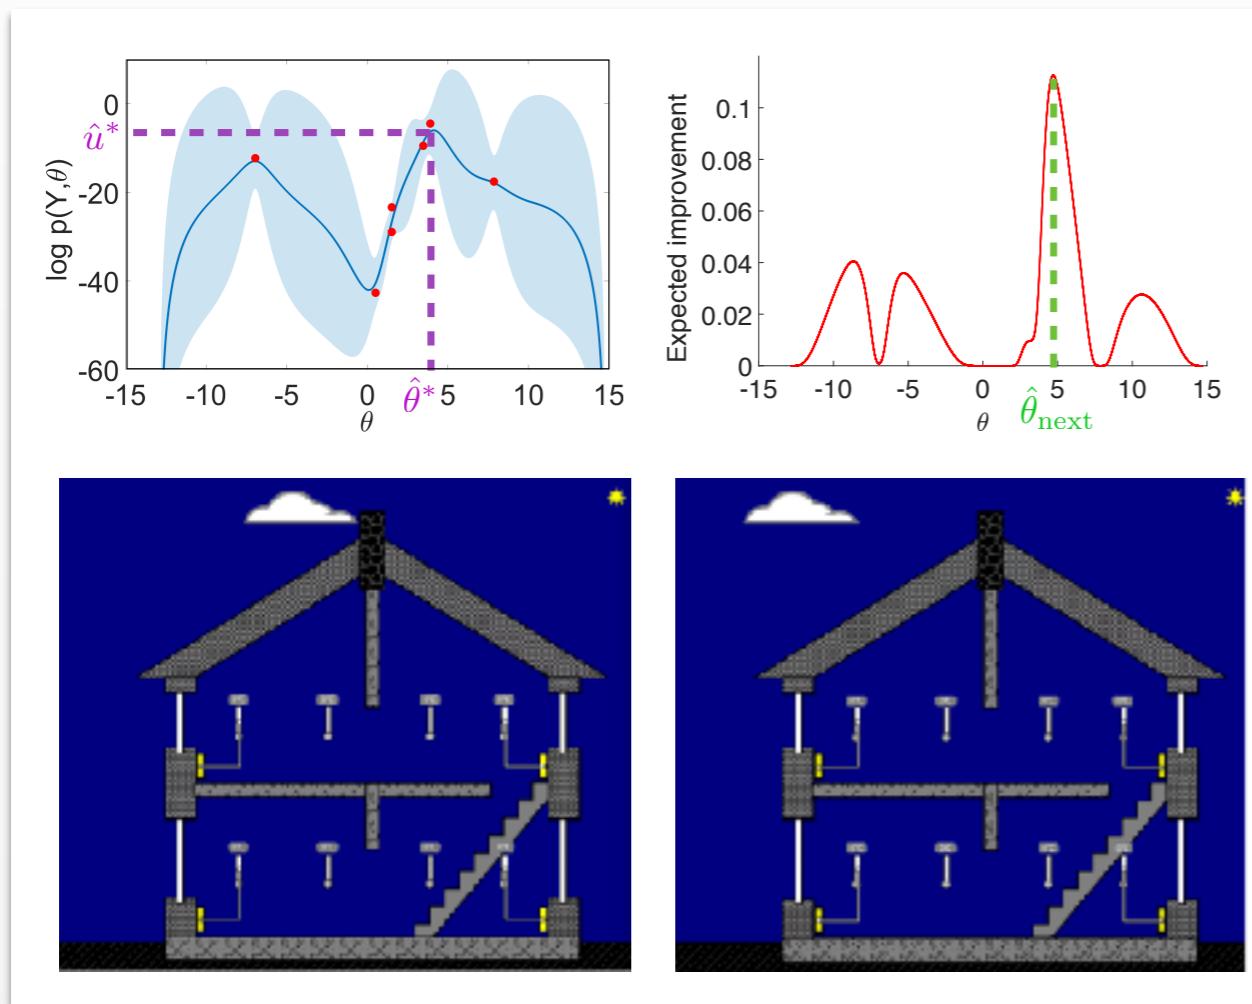


Babak
Esmaeili

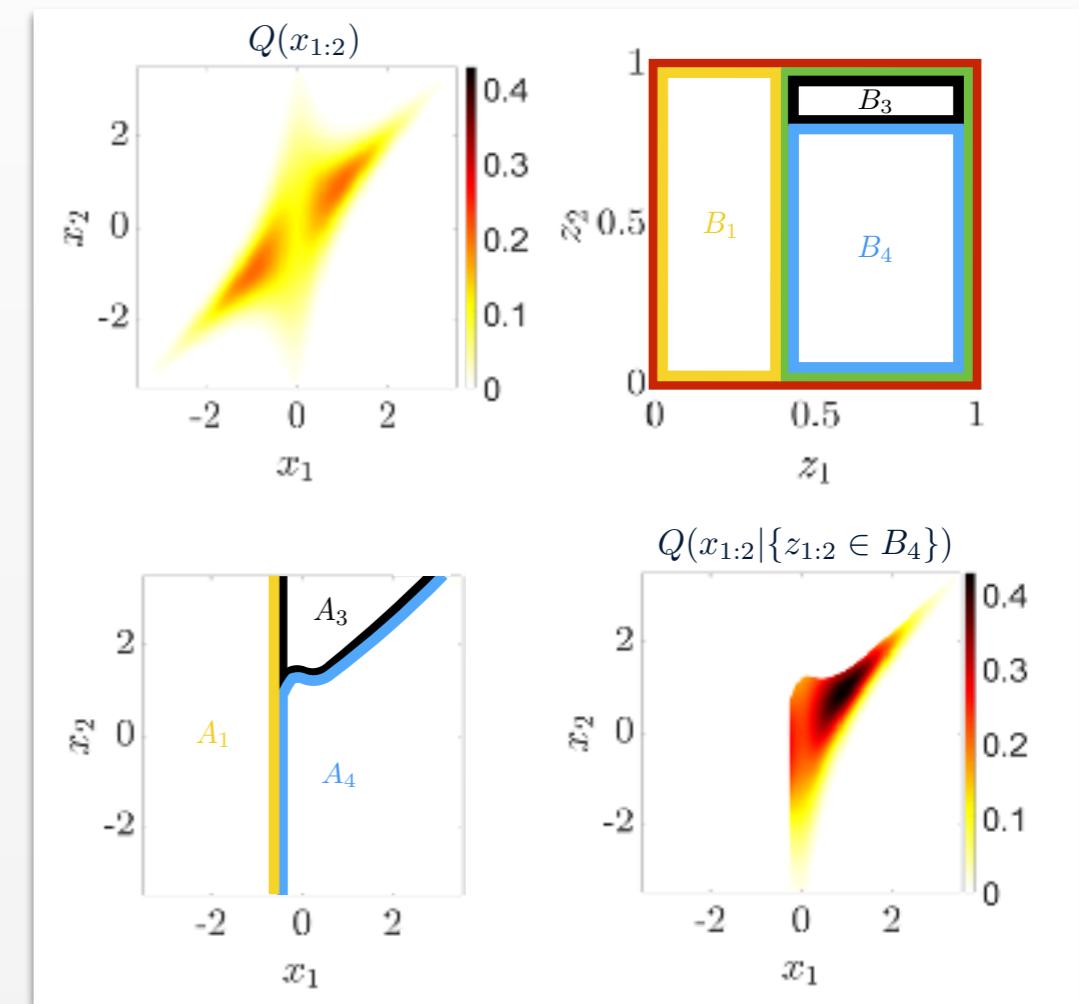
We're hiring: Faculty, Postdocs, Phd Students

Exploration vs Exploitation in Inference

Bayesian Optimization



Inference Trees

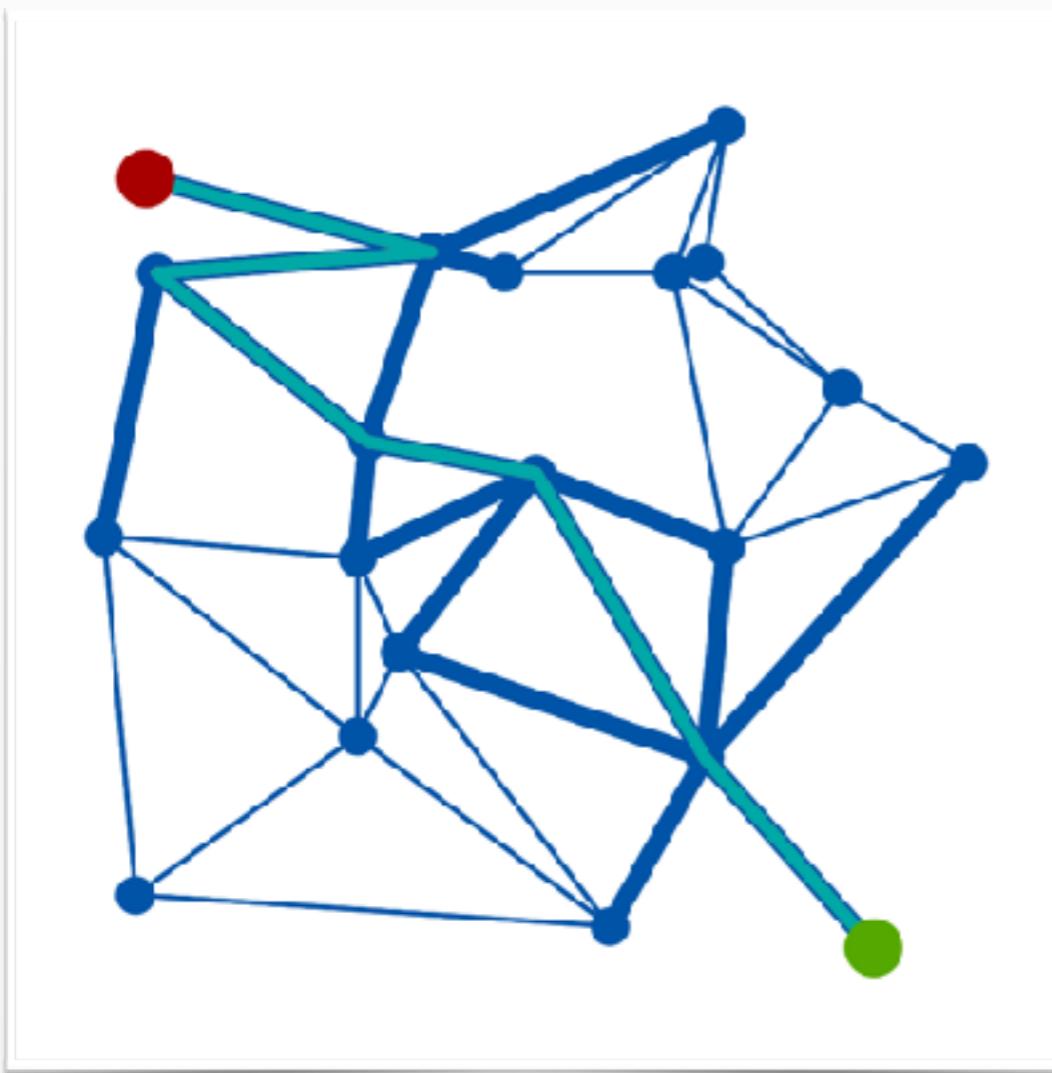


Idea: Use probabilistic programs to define marginal MAP problems
[Rainforth et al. 2016]

Idea: Partition space, use MCTS to allocate computation
[Rainforth et al. 2017]

Implementation Example

Canadian Traveler Problem

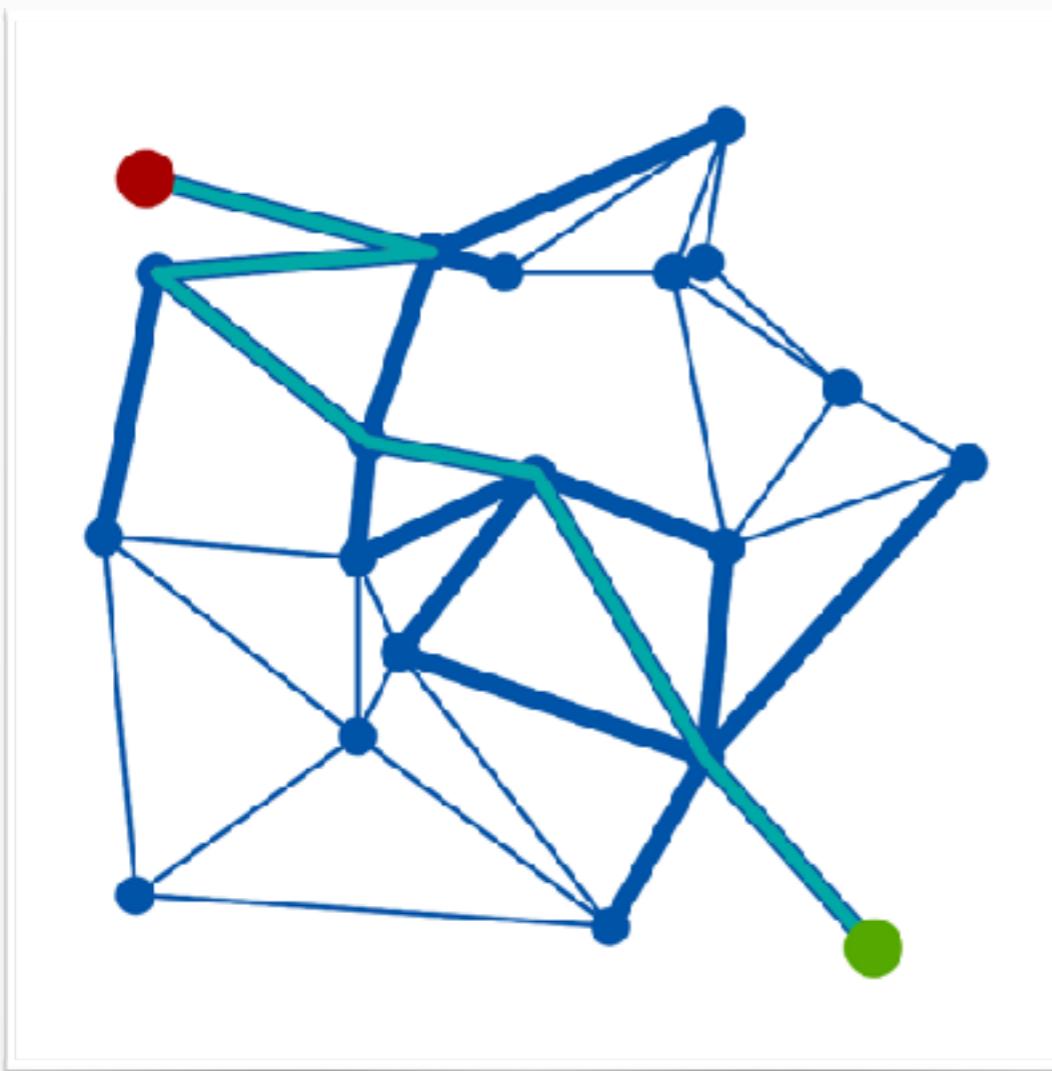


Model Specification

```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                  graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                  instance
                  start target
                  (make-policy)))]
    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))
    (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



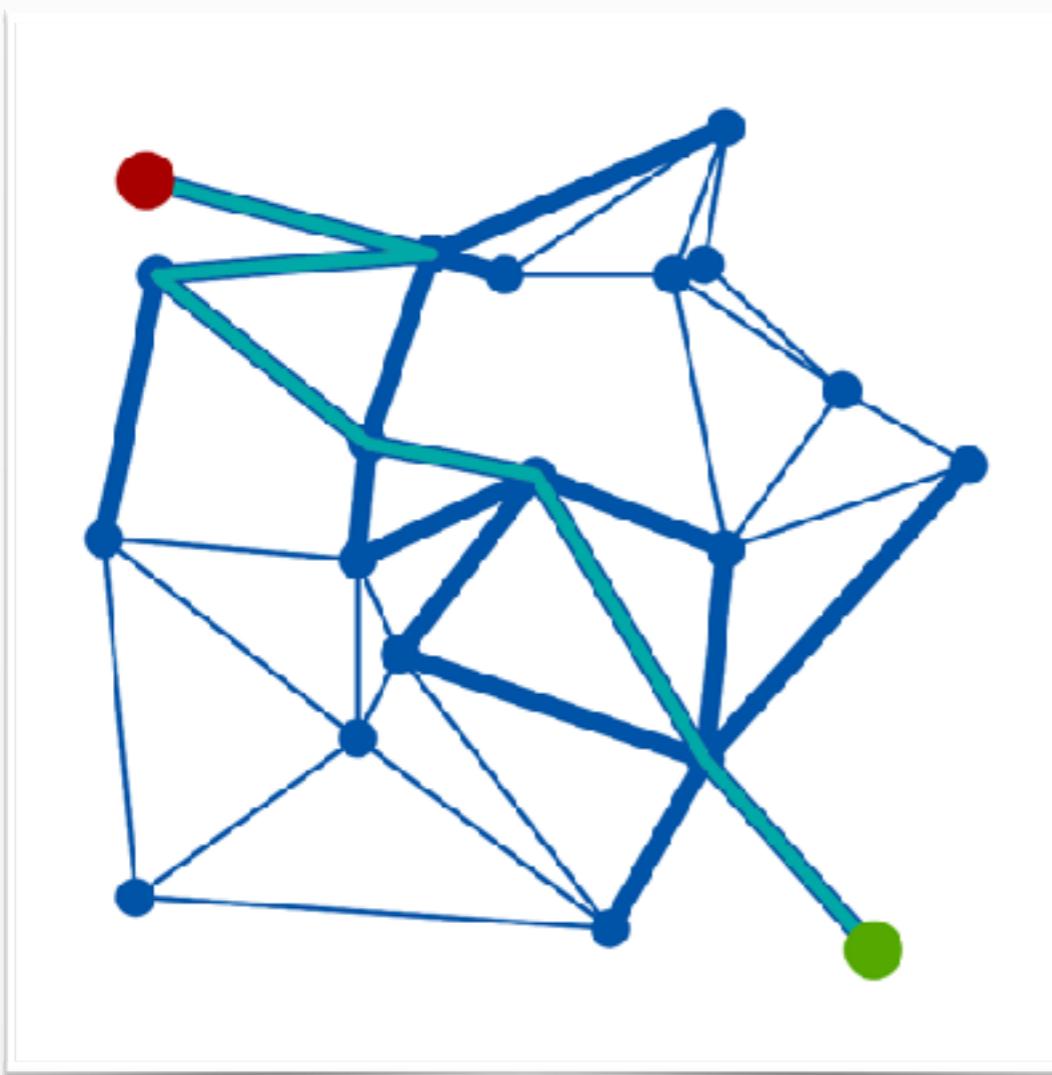
Model Specification

```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                  graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                  instance
                  start target
                  (make-policy))]

    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))
    (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



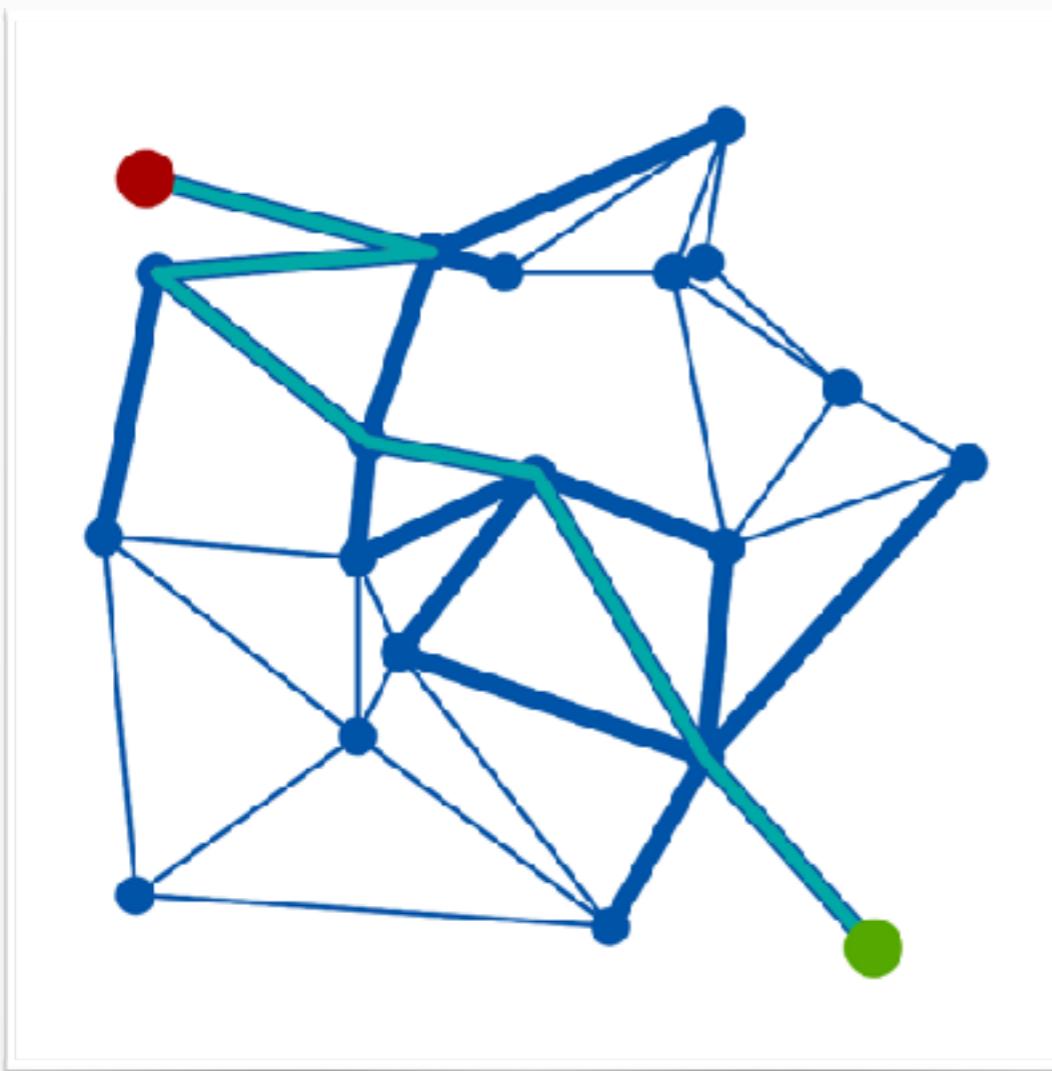
Model Specification

```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                    graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                   instance
                   start target
                   (make-policy))]

    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))
    (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



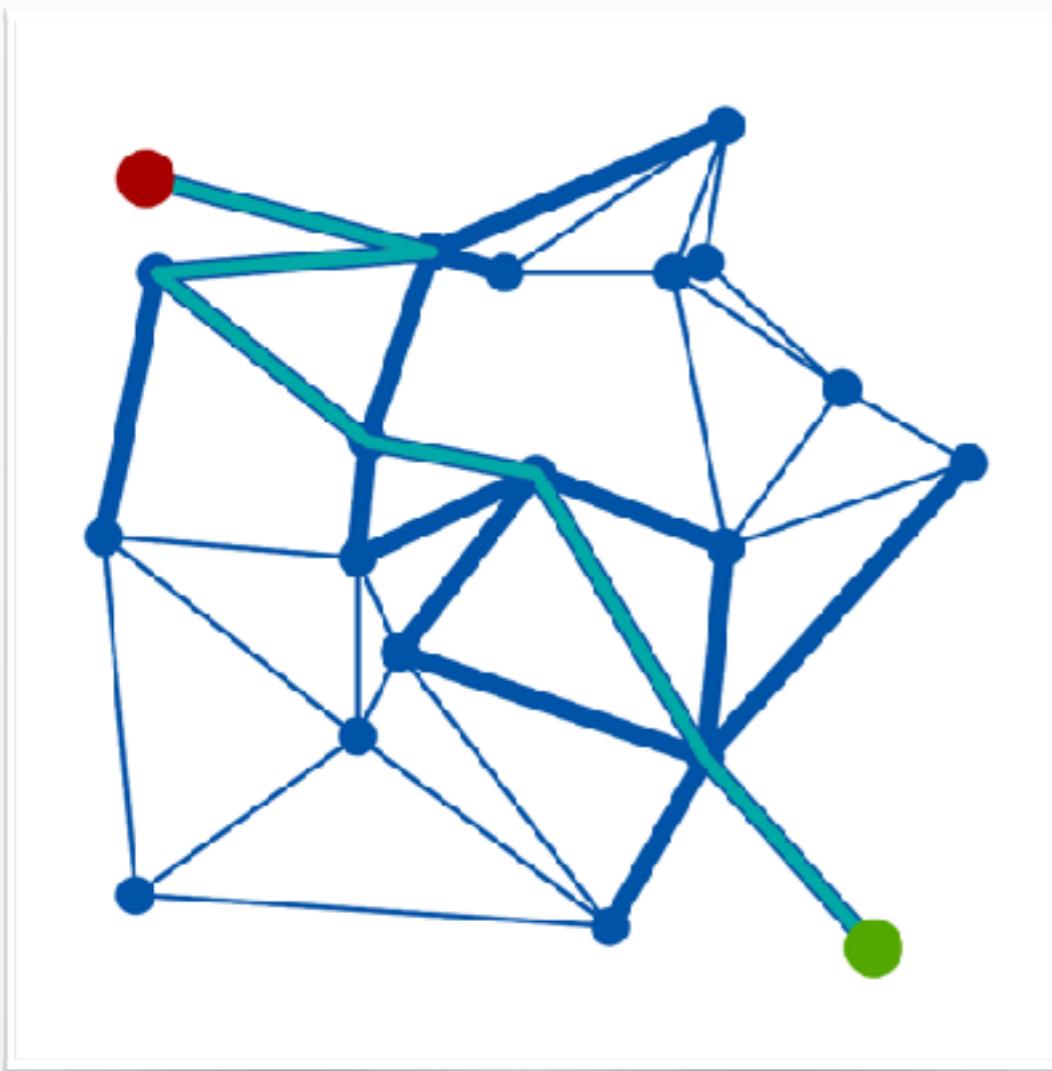
Model Specification

```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                  graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                  instance
                  start target
                  (make-policy))]

    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))
    (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



Model Specification

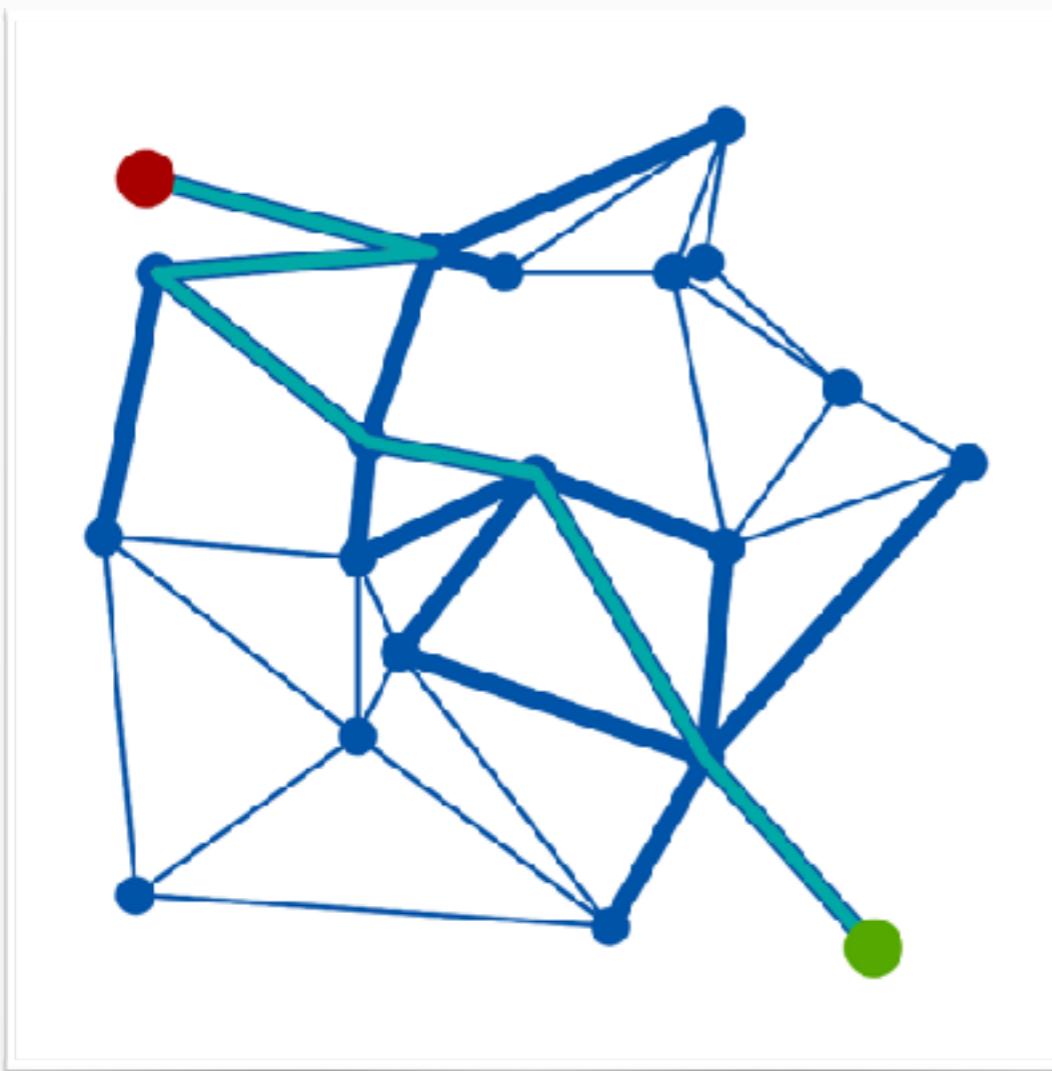
```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                  graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                  instance
                  start target
                  (make-policy))]

    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))))

  (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



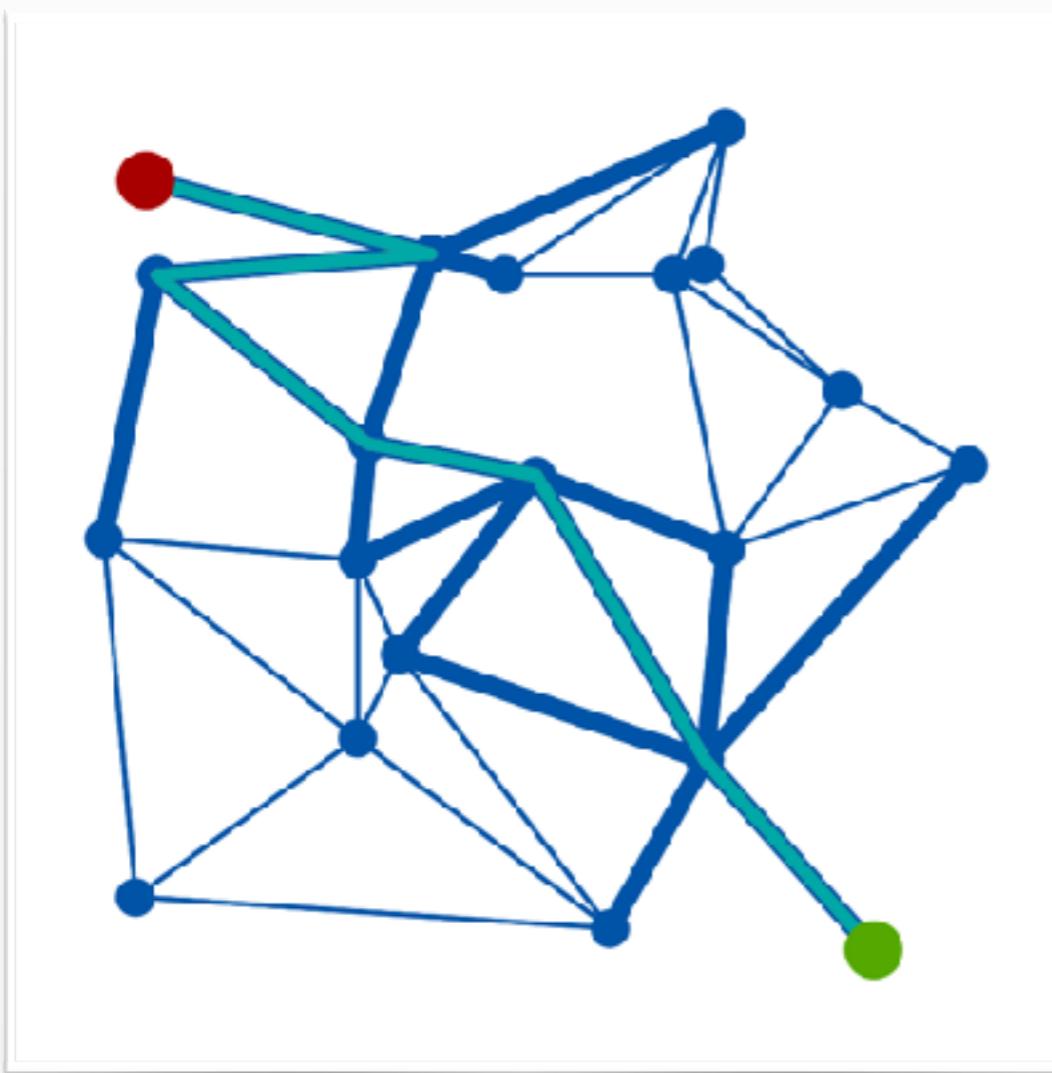
Model Specification

```
(defquery ctp
  [graph start target
   open-prob make-policy]
  (let [;; sample open/closed edges
        instance (simulate-weather
                  graph open-prob)
        ;; do depth-first search
        distance (dfs-agent
                  instance
                  start target
                  (make-policy))]

    ;; weigh according to distance,
    ;; reject when not connected
    (factor (- (or distance inf)))
    (predict :distance distance))))
```

Implementation Example

Canadian Traveler Problem



Model Specification

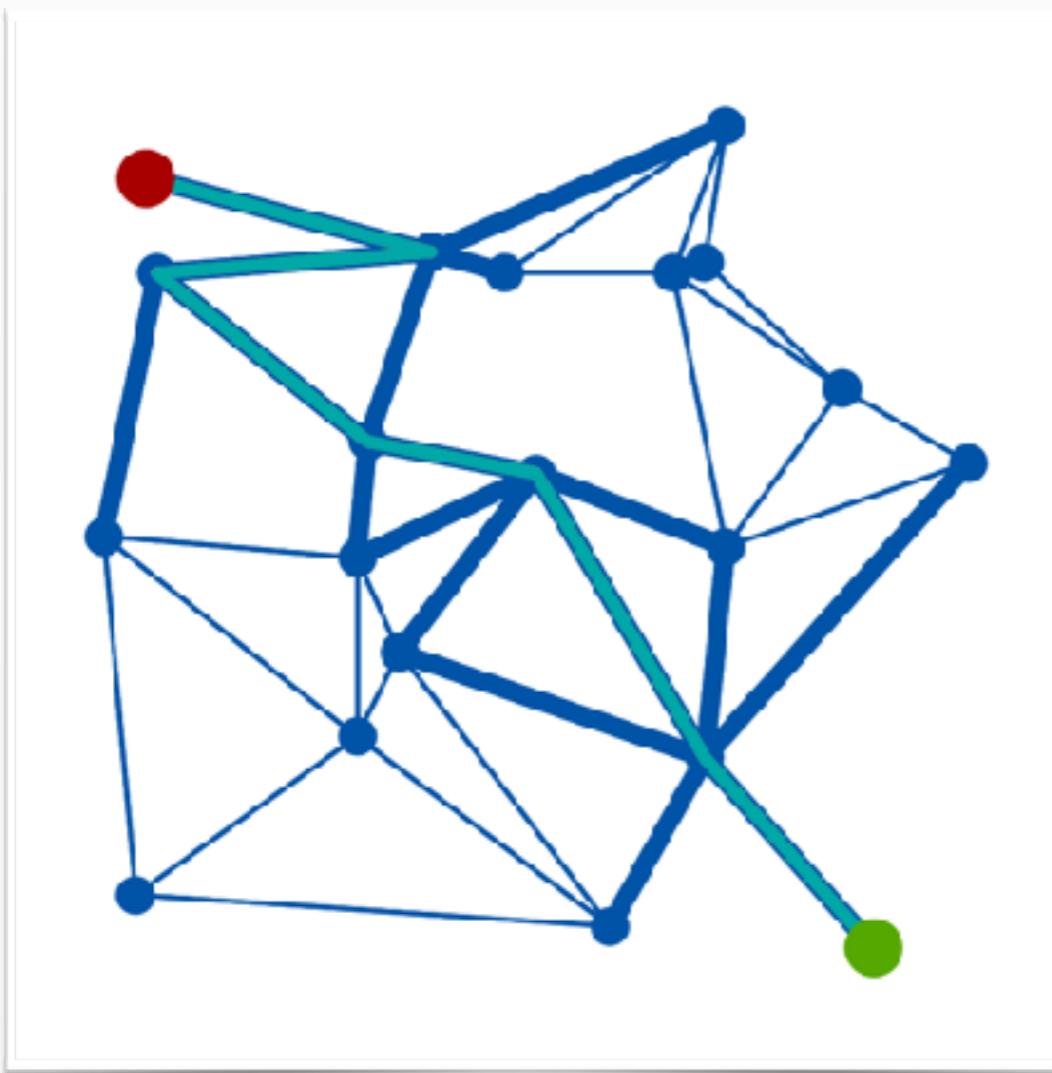
```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
                  (sample
                    (learn [u v]
                      (gamma 1. 1.))))))
        (fn policy [u opts]
          (let [Qs (map (fn [v]
                          (Q u v)) opts)]
            (argmax (zipmap opts Qs))))])
    (fn policy [u opts]
      (let [Qs (map (fn [v]
                      (Q u v)) opts)]
        (argmax (zipmap opts Qs)))))))
```

Implementation Example

Canadian Traveler Problem



Model Specification

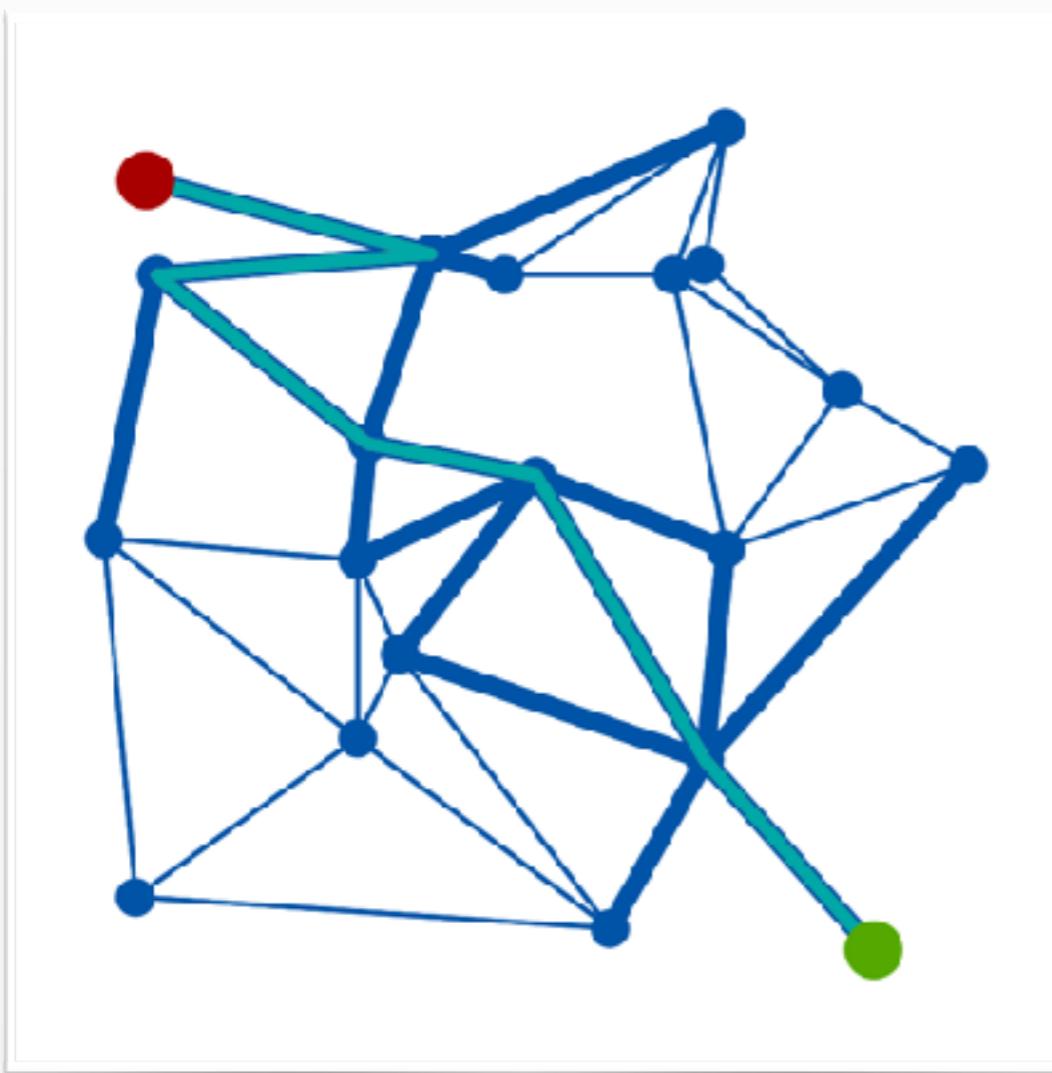
```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
                  (sample
                    (learn [u v]
                      (gamma 1. 1.))))))
        (fn policy [u opts]
          (let [Qs (map (fn [v]
                          (Q u v)) opts)]
            (argmax (zipmap opts Qs))))]))
```

Implementation Example

Canadian Traveler Problem



Model Specification

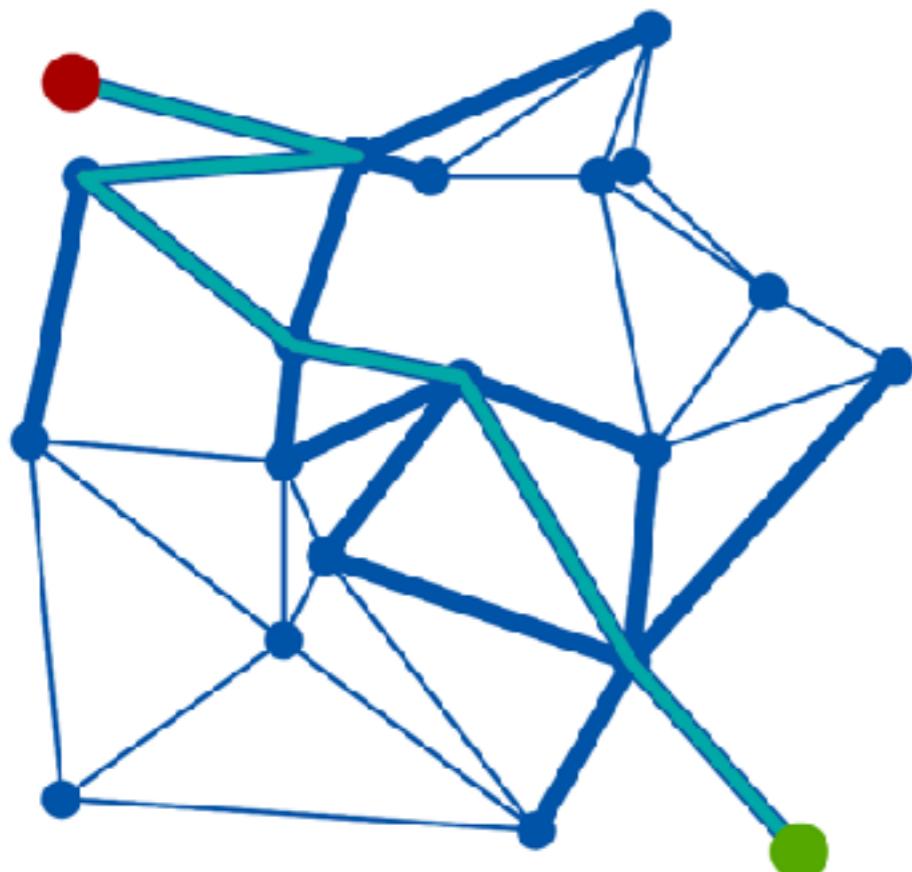
```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
                  (sample
                    (learn [u v]
                      (gamma 1. 1.))))))
        (fn policy [u opts]
          (let [Qs (map (fn [v]
                          (Q u v)) opts)]
            (argmax (zipmap opts Qs))))])
    (fn policy [u opts]
      (let [Qs (map (fn [v]
                      (Q u v)) opts)]
        (argmax (zipmap opts Qs)))))))
```

Implementation Example

Canadian Traveler Problem



Model Specification

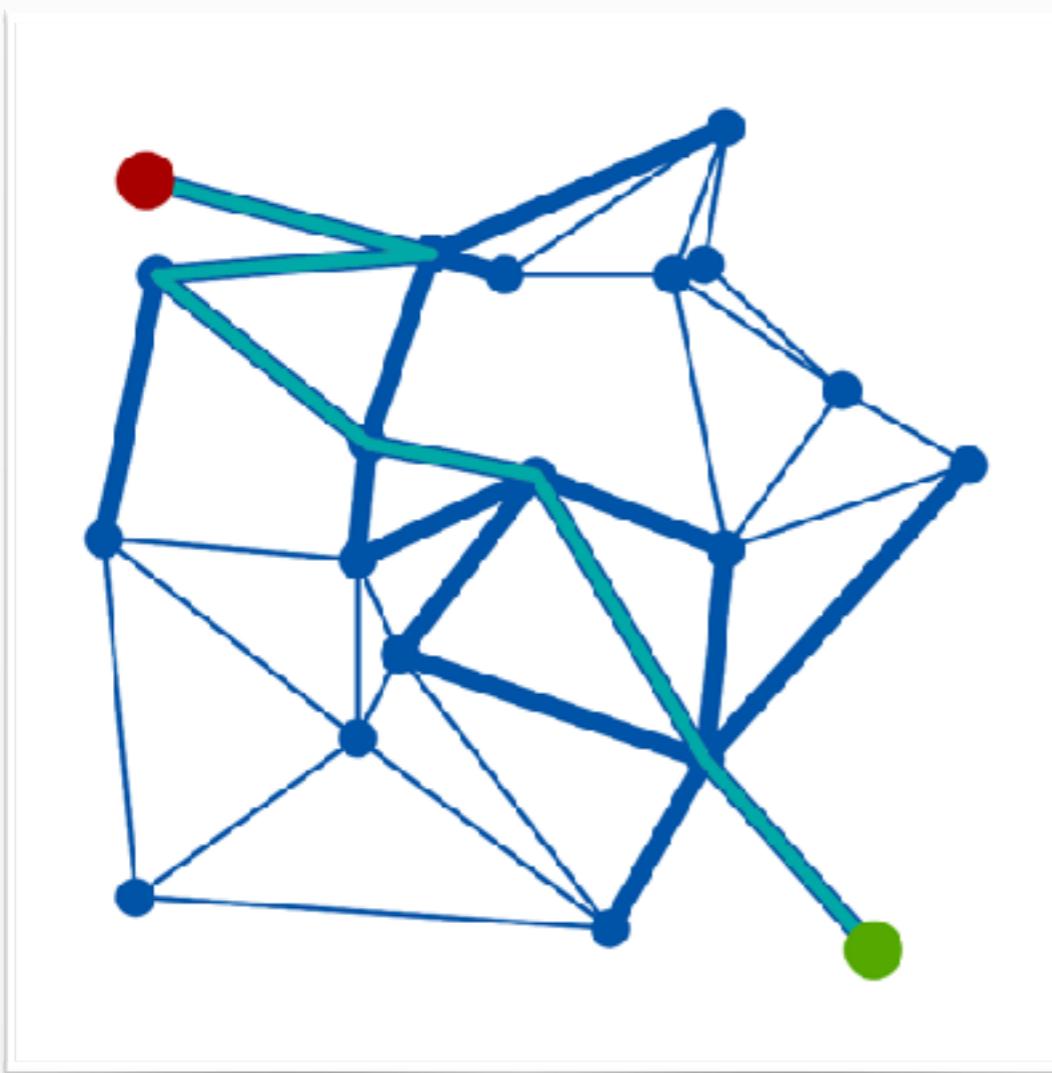
```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
                  (sample
                    (learn [u v]
                      (gamma 1. 1.)))))])
    (fn policy [u opts]
      (let [Qs (map (fn [v]
                      (Q u v)) opts)]
        (argmax (zipmap opts Qs))))))
```

Implementation Example

Canadian Traveler Problem



Model Specification

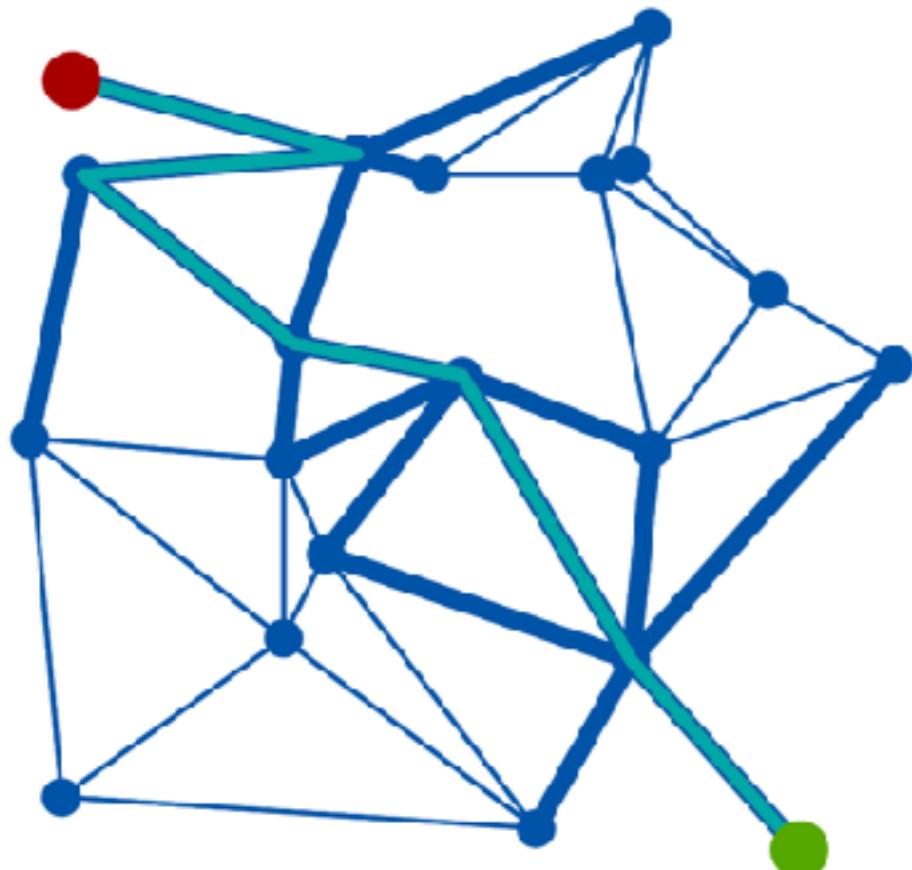
```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
                  (sample
                    (learn [u v]
                      (gamma 1. 1.)))))])
    (fn policy [u opts]
      (let [Qs (map (fn [v]
                      (Q u v)) opts)]
        (argmax (zipmap opts Qs))))))
```

Implementation Example

Canadian Traveler Problem



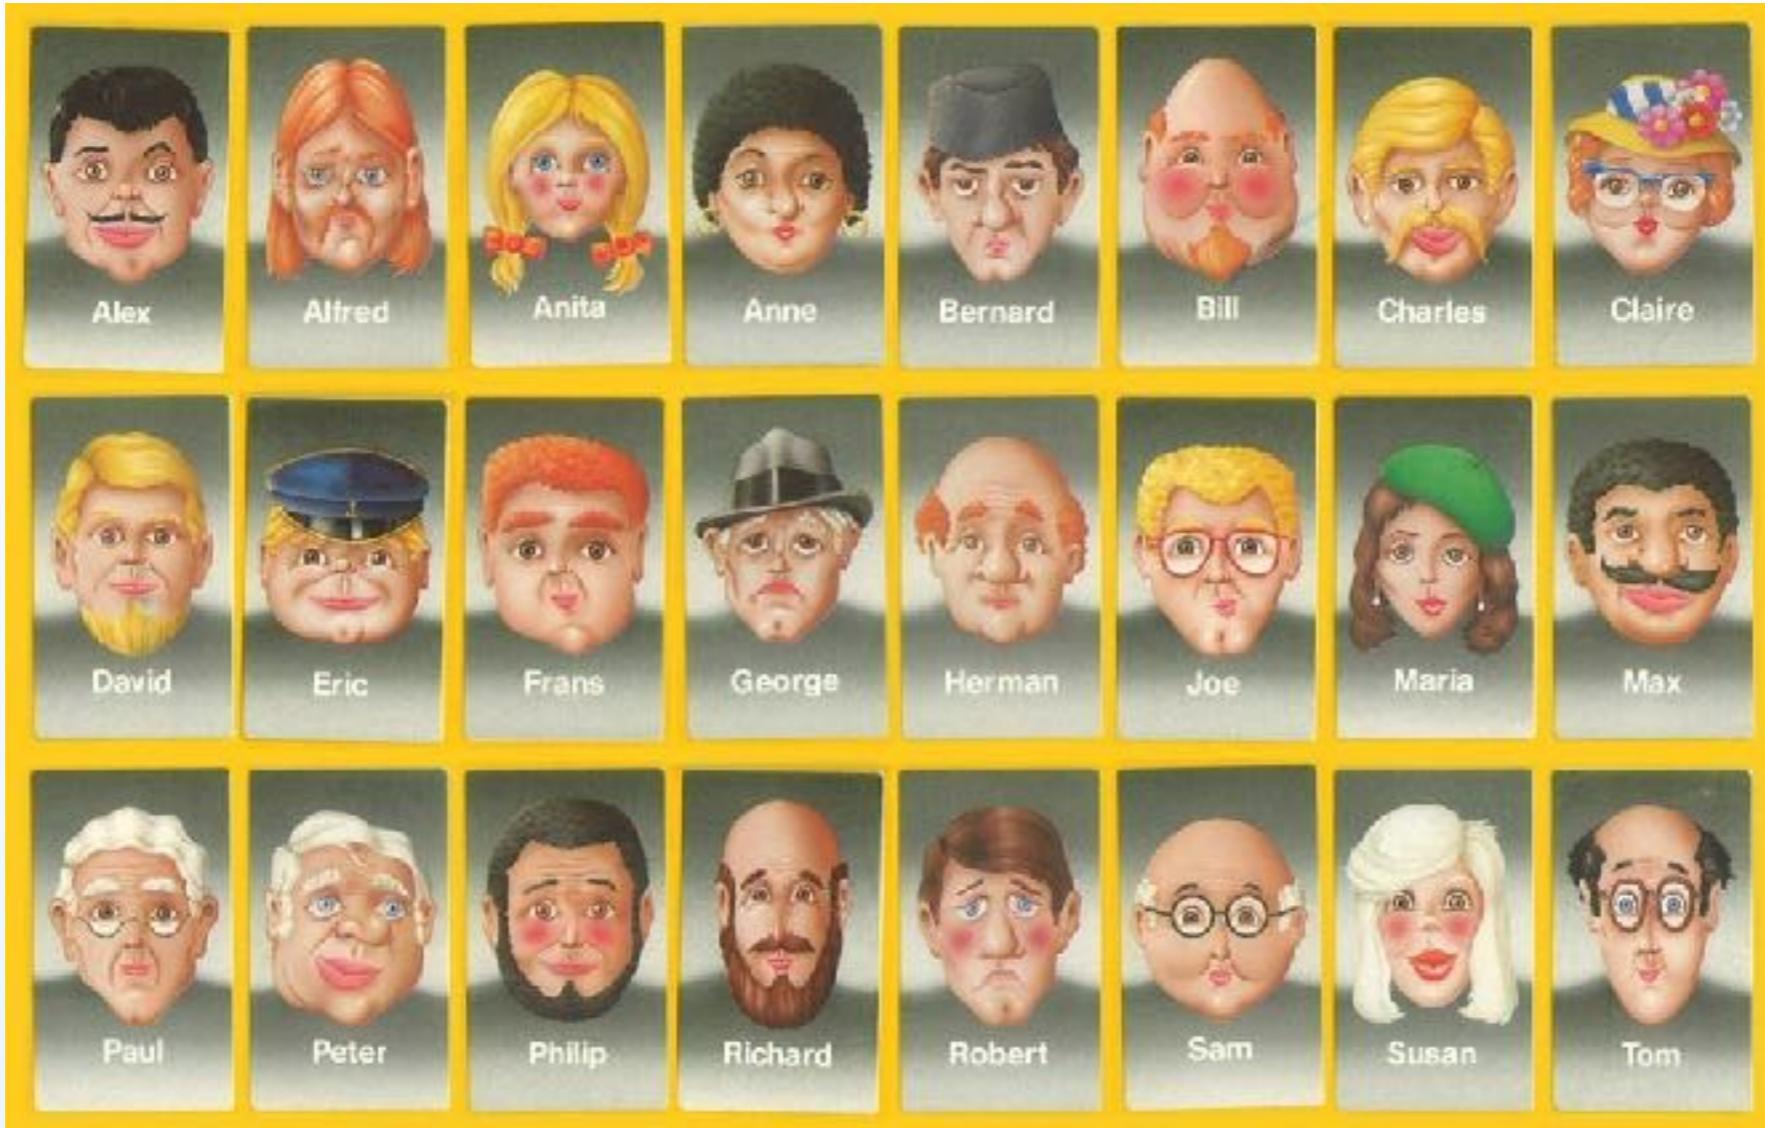
Model Specification

```
(defm make-random-policy []
  (fn policy [u opts]
    (sample (uniform opts))))
```

Utility-based Policy

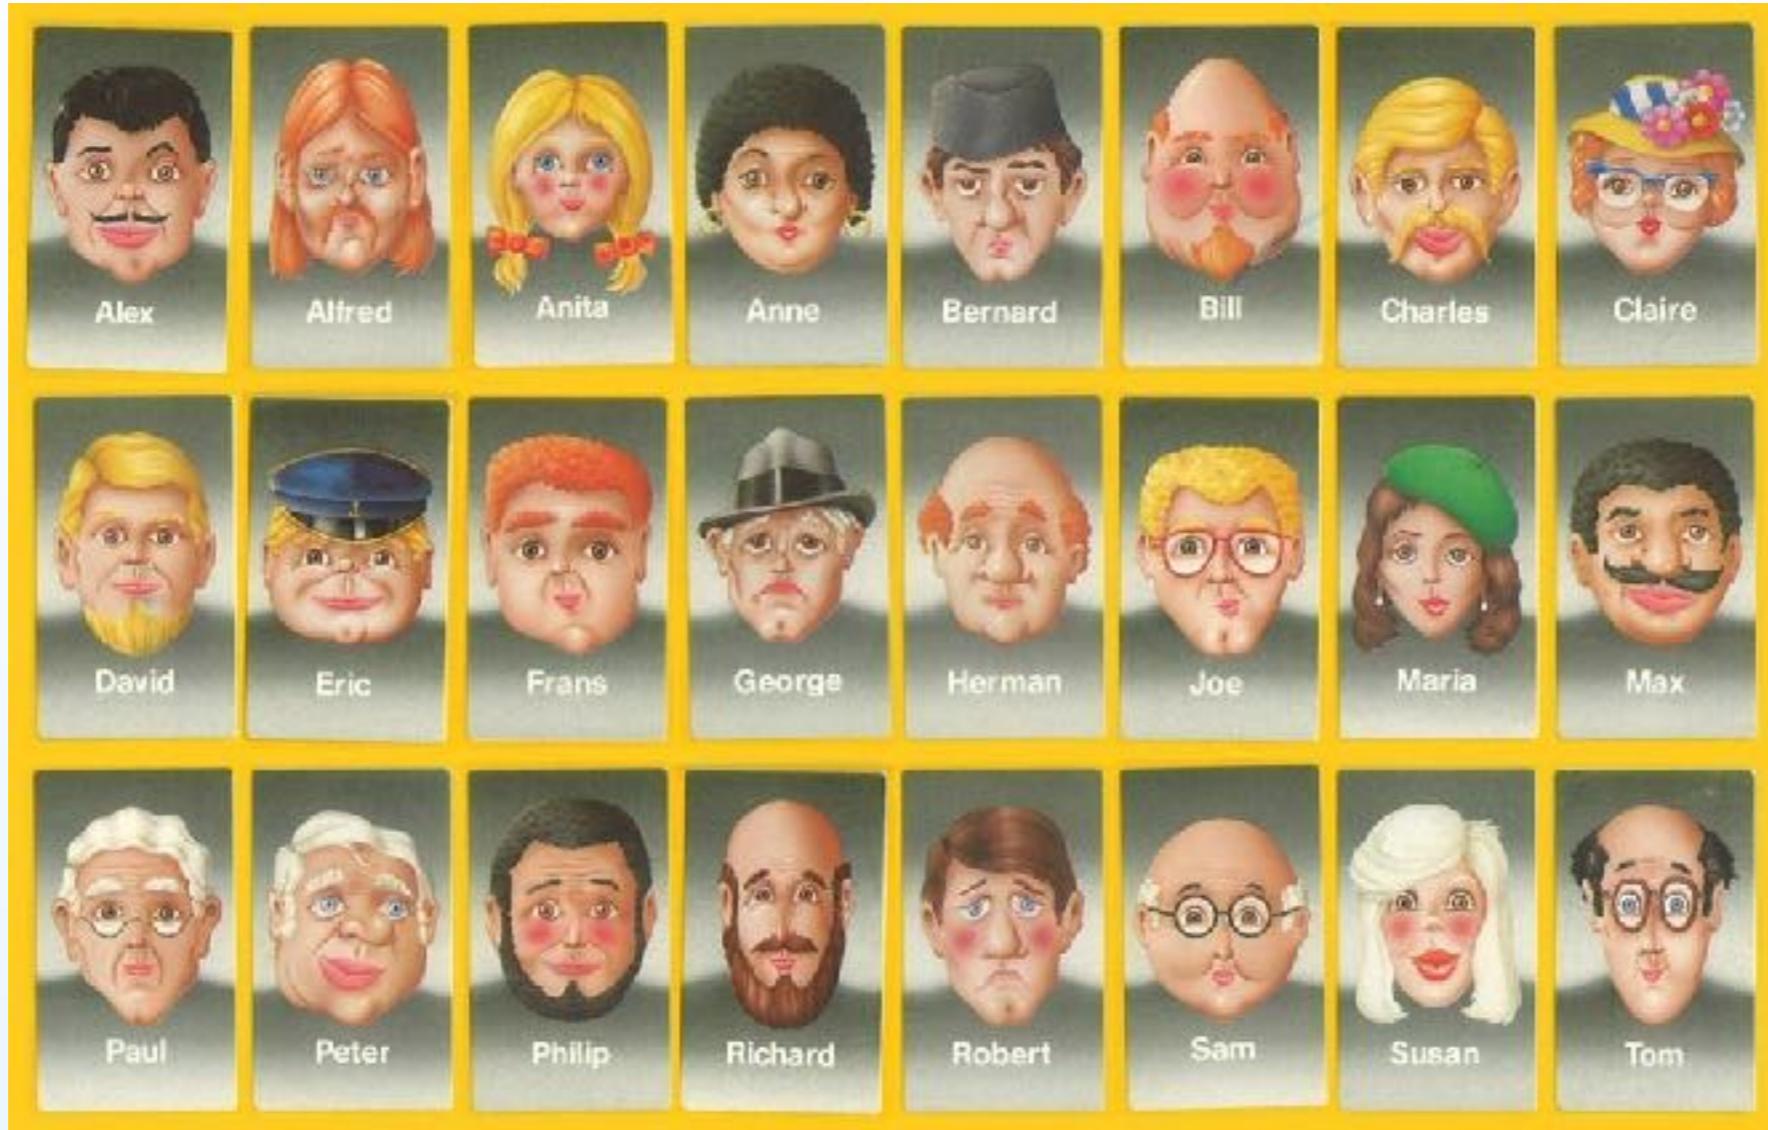
```
(defm make-edge-policy []
  (let [Q (mem (fn [u v]
      (sample
        (learn [u v]
          (gamma 1. 1.)))))])
    (fn policy [u opts]
      (let [Qs (map (fn [v]
          (Q u v)) opts)]
        (argmax (zipmap opts Qs))))))
```

Example: Guess Who



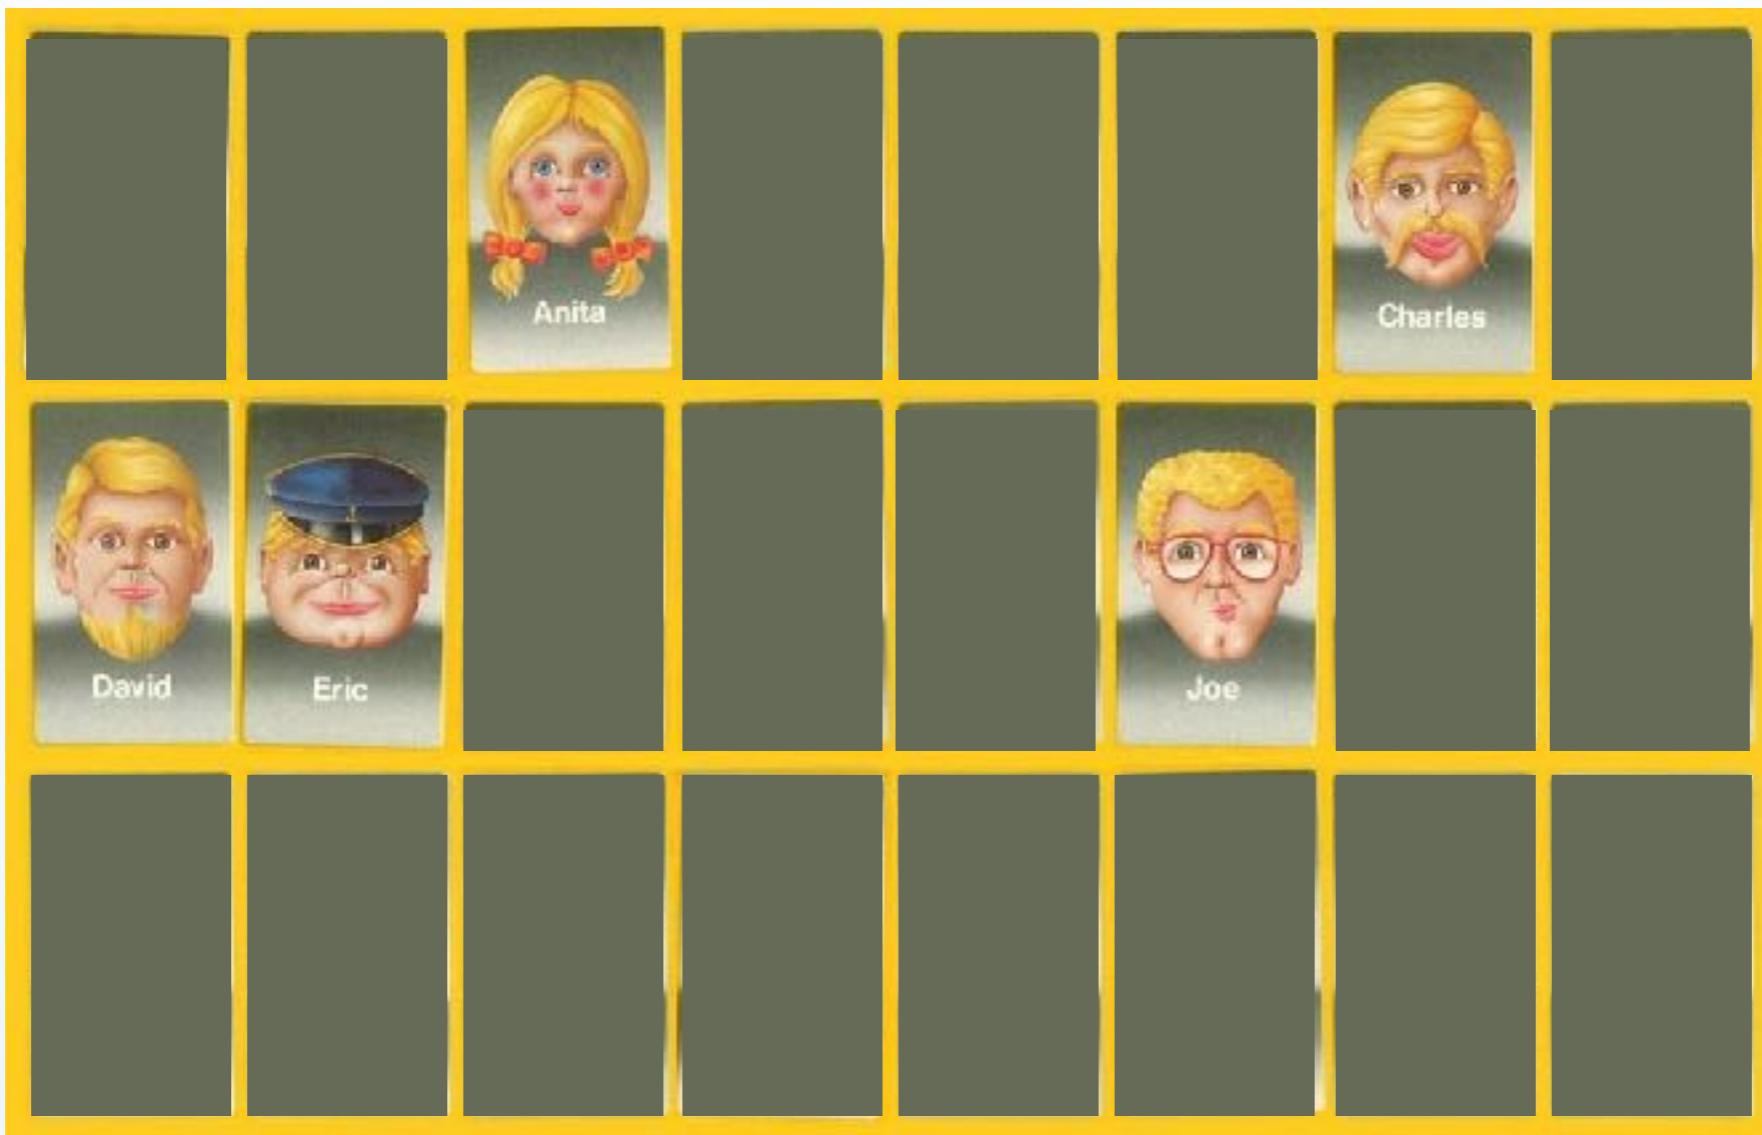
Premise: Player asks questions
to identify an individual in a set

Example: Guess Who



Question: Does the person have **blonde hair**?

Example: Guess Who



Question: Does the person have **blonde hair?**

Yes: 5 candidates, No: 19 candidates

Example: Guess Who



Question: Does the person have **blonde hair?**

Yes: 5 candidates, **No: 19 candidates**

Example: Guess Who



Question: Does the person have **short hair**?

Yes: 15 candidates, No: 9 candidates

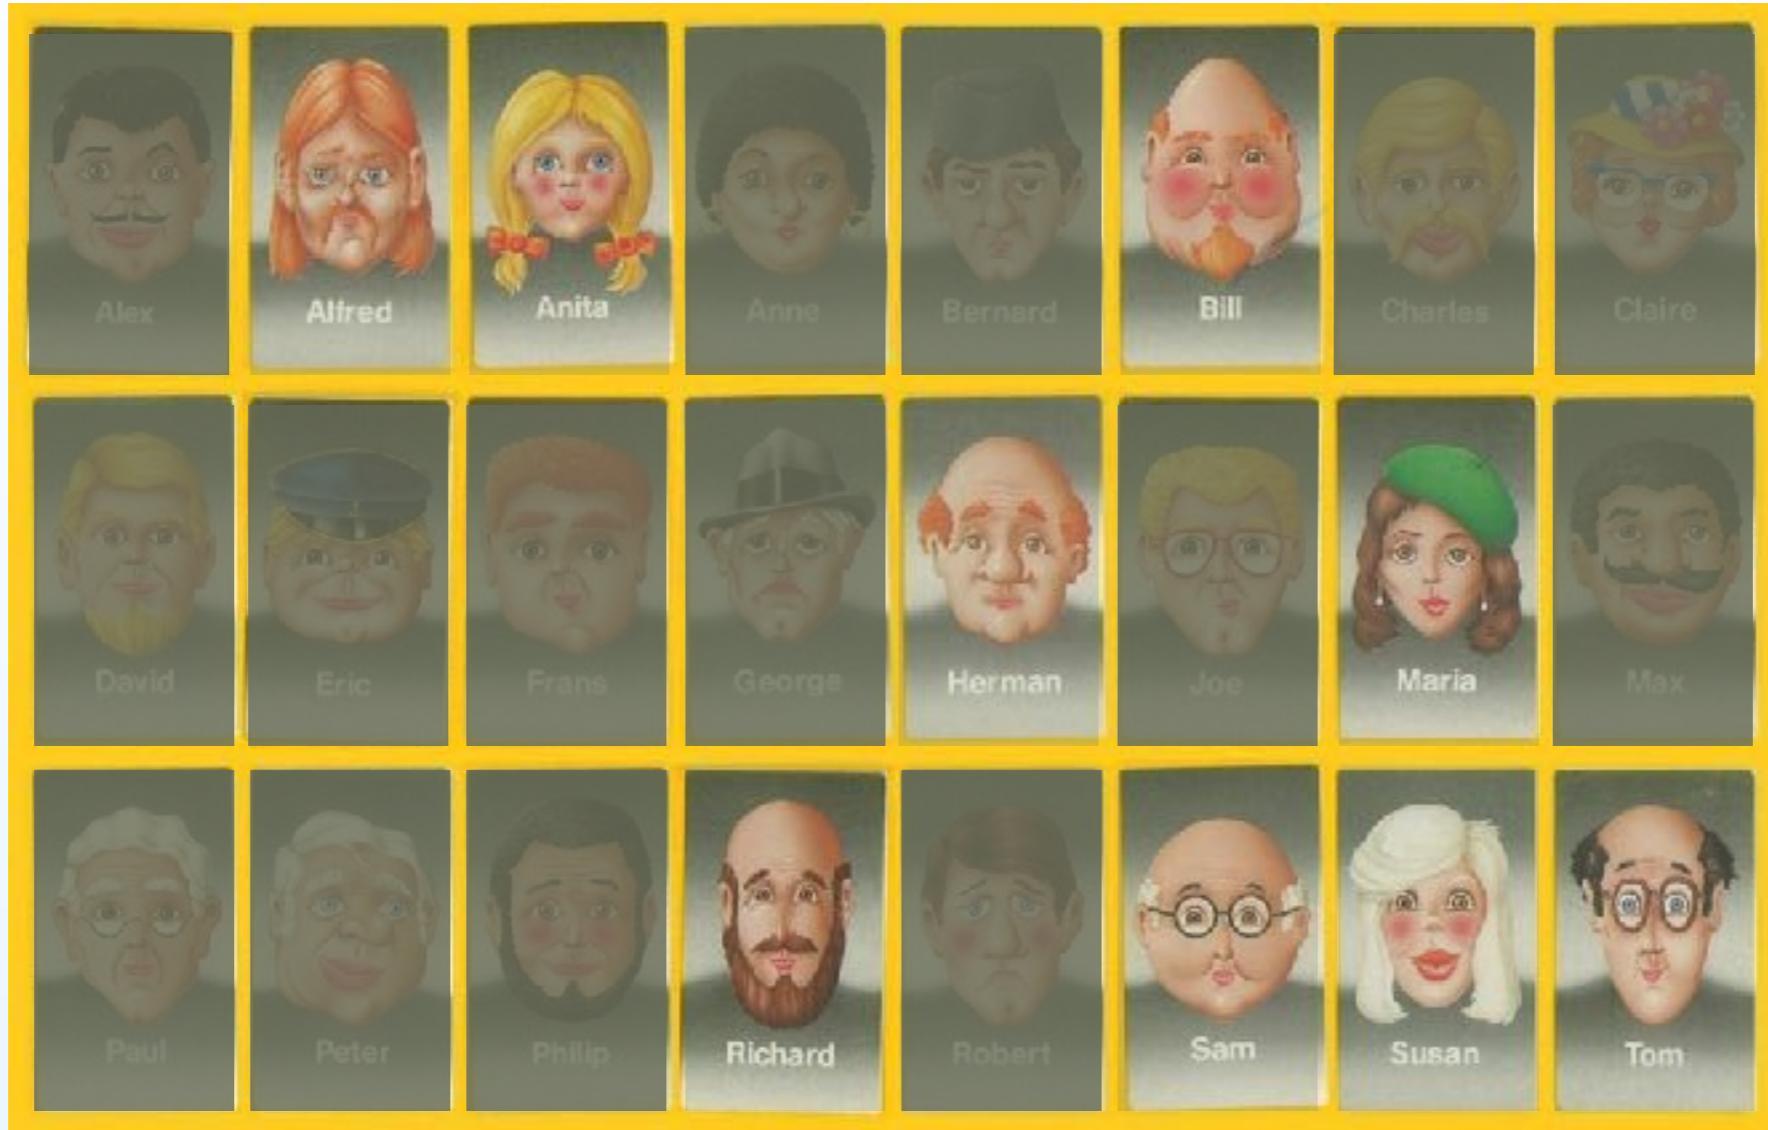
Example: Guess Who



Question: Does the person have **blonde hair?**

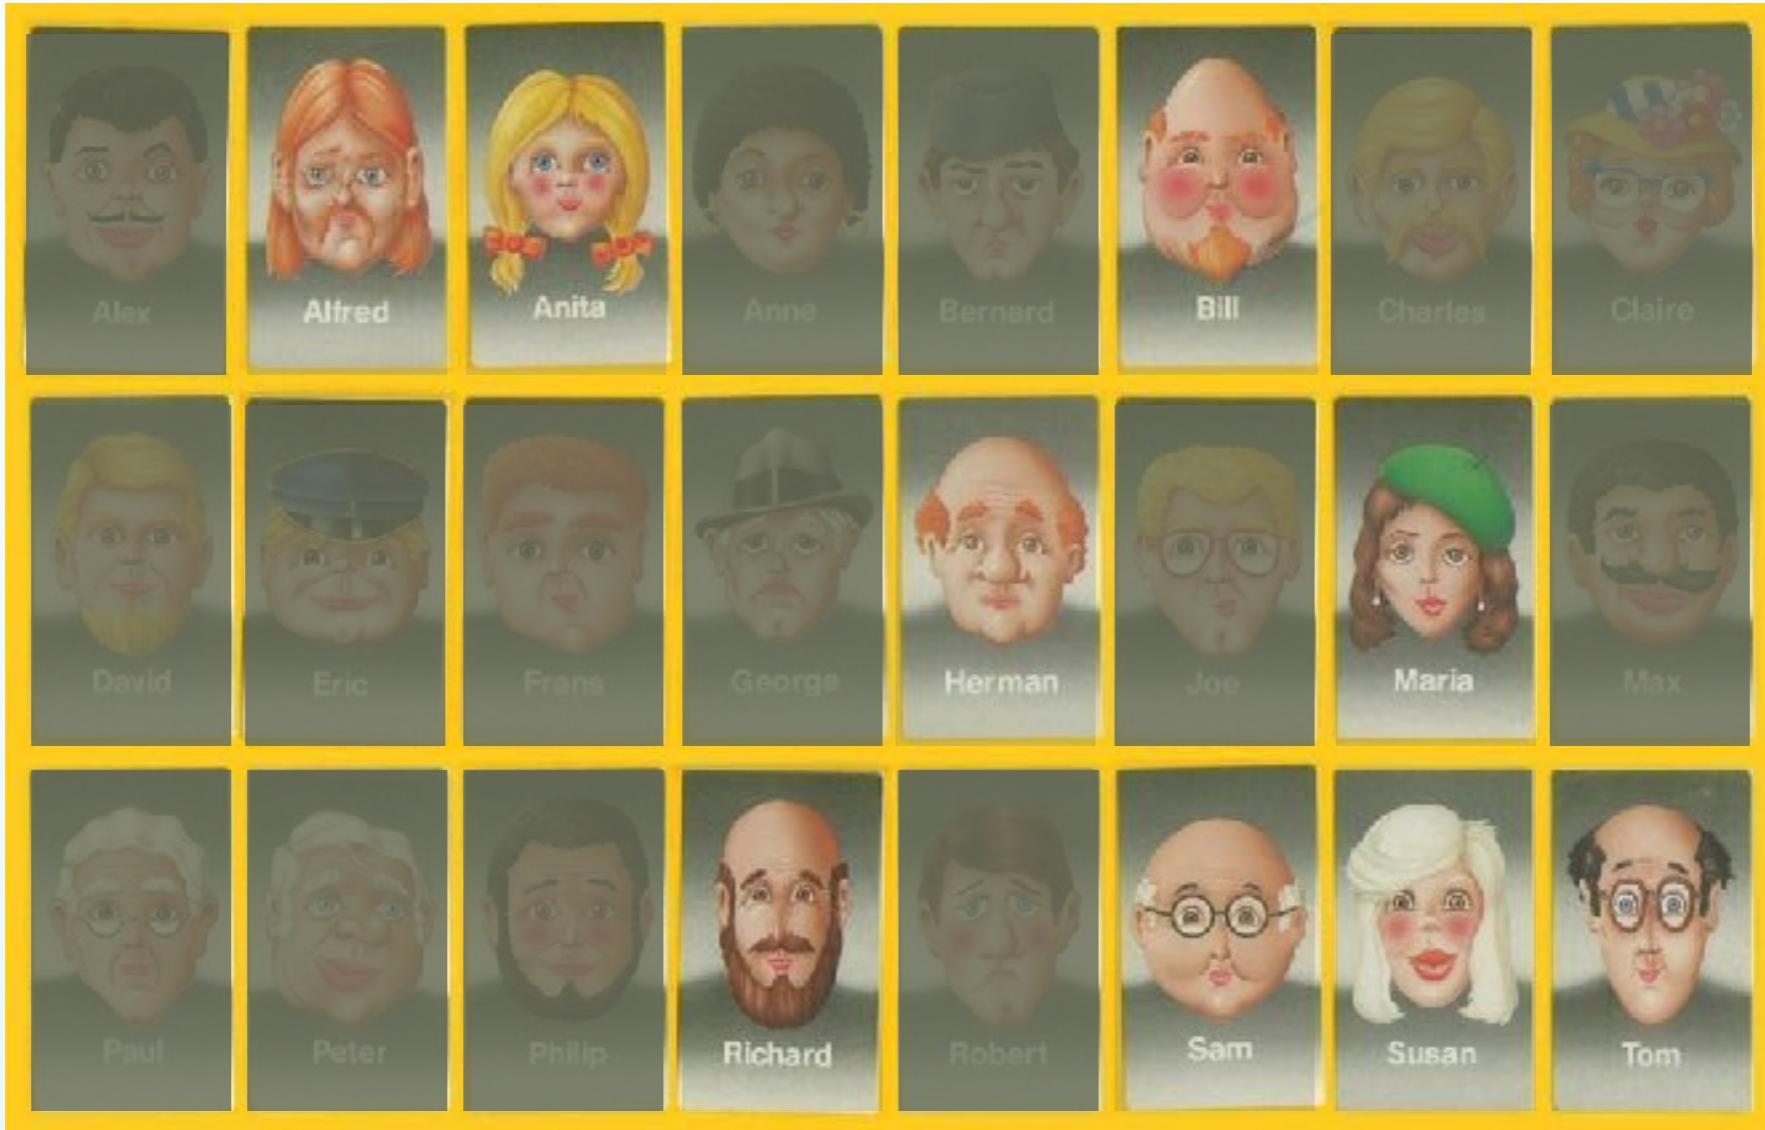
Yes: 15 candidates, **No: 9 candidates**

Example: Guess Who



What if answers are only 80% reliable?

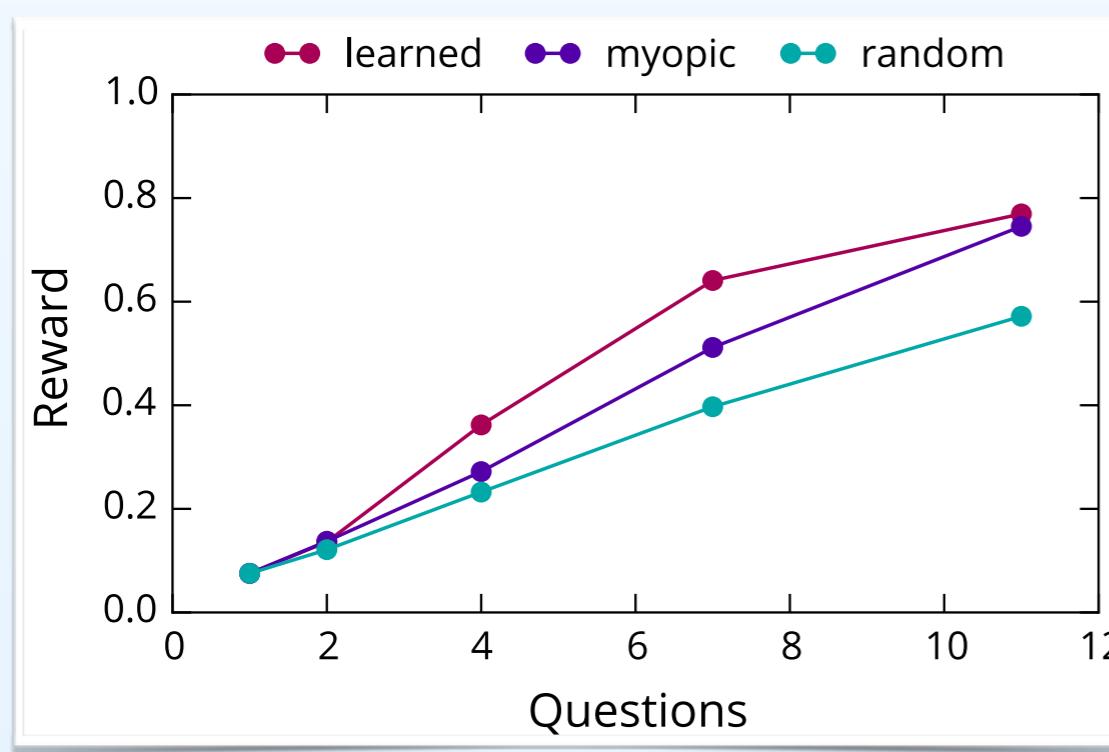
Example: Guess Who



What if answers are only 80% reliable?

What question ***is most likely to inform*** my current belief?

Guess Who (noisy)



Heuristic: Value of Information

- Difference between expected return for *current* best guess that for *new* best guess.
- Make myopic approximation (assume question is last in game)

Policy: Discounted Utility

$$U_q = \gamma^{c_q} (A \cdot b)_q$$

- U_q Utility of question q
 γ Heuristic discount factor
 c_q Question count
 A Weight matrix
 b Belief vector